

# Operating Systems

Timo Hönig

Bochum Operating Systems and System Software (BOSS)

Ruhr University Bochum (RUB)

V. Deadlocks

May 10, 2023 (Summer Term 2023)



RUHR  
UNIVERSITÄT  
BOCHUM

RUB

[www.informatik.rub.de](http://www.informatik.rub.de)

Chair of Operating Systems and System Software



- **process synchronisation**

- **critical sections, race conditions**
- coordination methods to **avoid errors** due to uncontrolled, **concurrent data access**

- ad-hoc methods using **busy waiting**

- **inefficient** solutions, dangerous illusion of correctness
- **waste of resource** (i.e., CPU cycles)

- advanced methods to ensure mutual exclusion:  
**passive waiting**

- exploit hardware support: **atomic operations**
- operating system support: **semaphores** to implement **passive waiting**





# Organizational Matters

- lecture
  - Wednesday, 10:15 – 11:45
  - format: synchronous, **hybrid**
    - in presence (Room HID, Building ID)
    - online lecture (Zoom)
  - **exam:** August 7, 2023 (first appointment)  
September 25, 2023 (retest appointment)
- exercises: **group allocation almost complete**
  - make use of group work - for your own benefit!
  - remarks on Wednesday for Übung T03
- manage course material, asynchronous communication: Moodle
- <https://moodle.ruhr-uni-bochum.de/course/view.php?id=50698>



# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ **Problem Scenario and Cause Study**
- ▶ Process Deadlocks
  - ▶ Definition, Variants, and Conditions
  - ▶ Consumable and Non-consumable Resources
  - ▶ Resource Allocation Graphs
- ▶ Classic Deadlock Situation
  - ▶ The Dining Philosophers Problem
- ▶ Deadlock Counter Measures
  - ▶ Prevention, Avoidance, Detection and Recovery
- Summary and Outlook

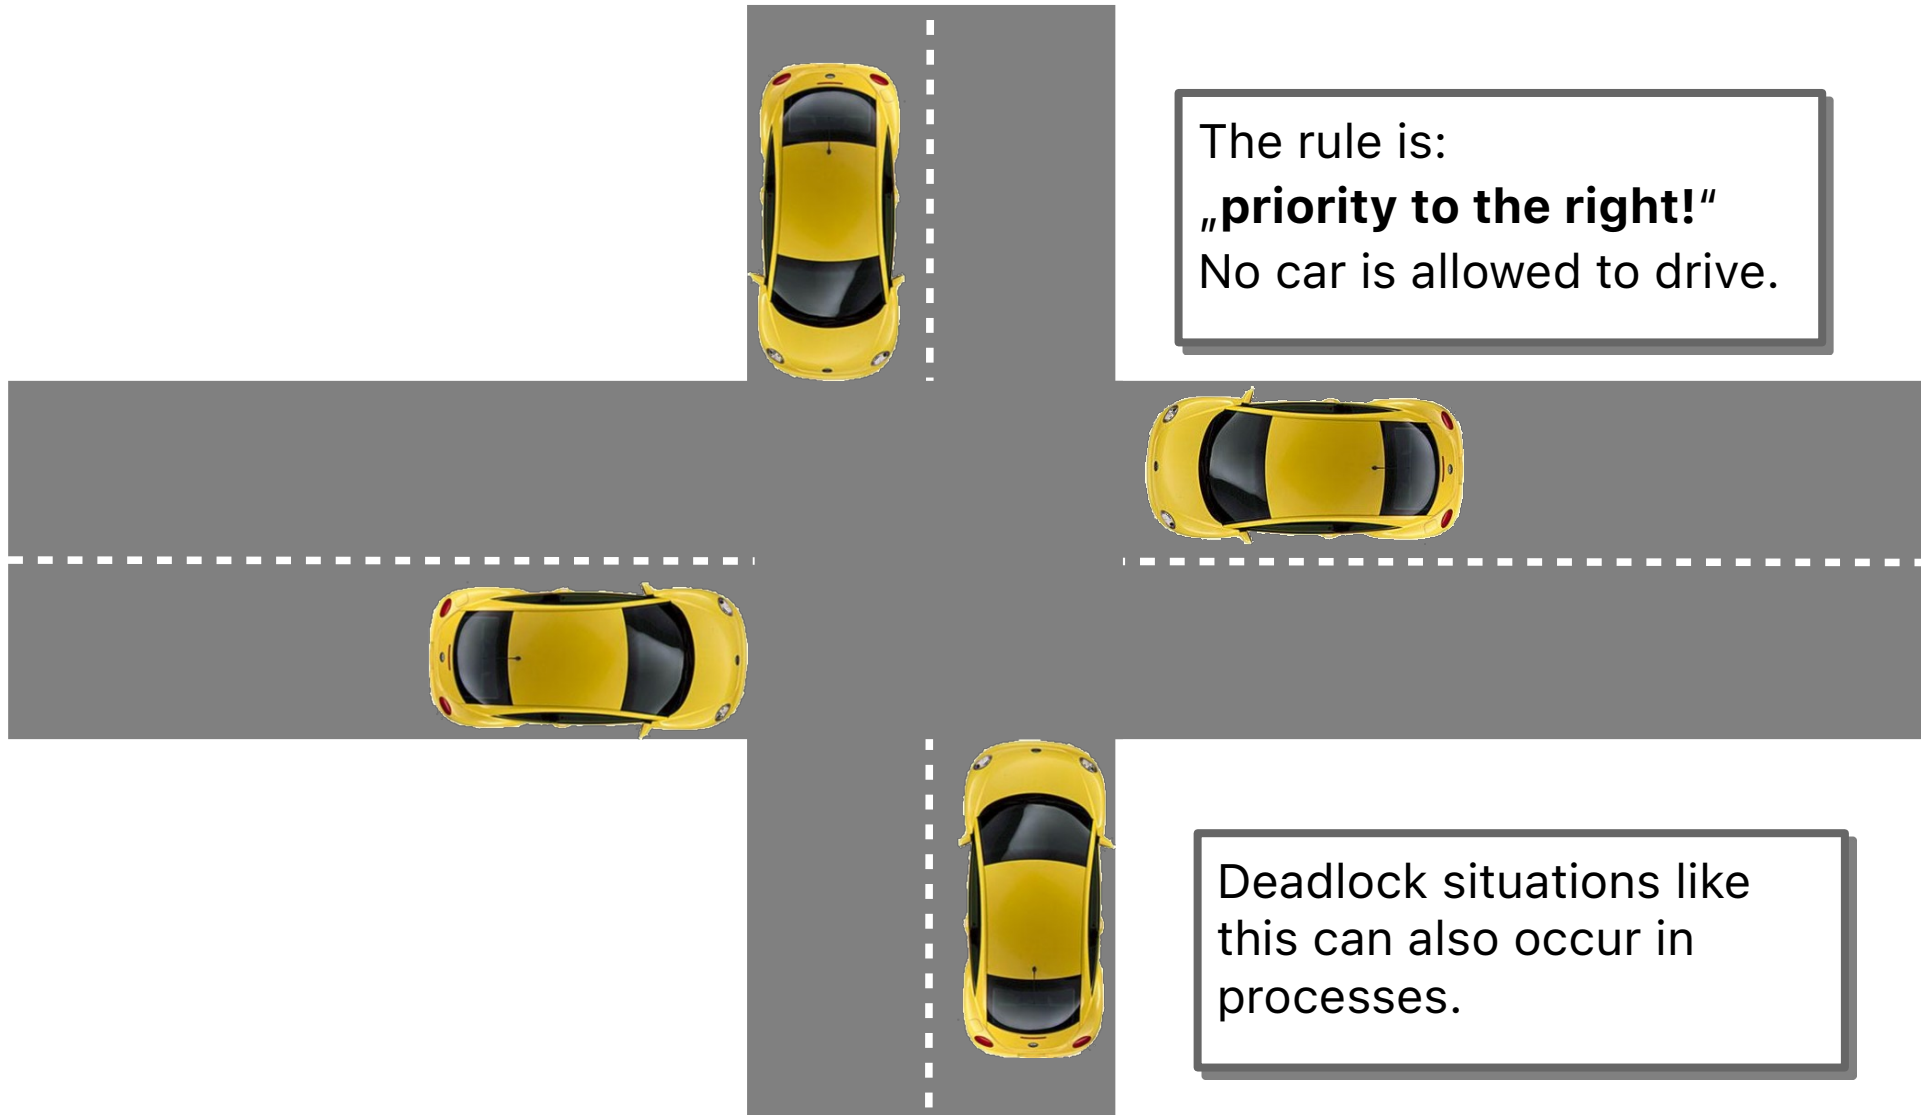


# Problem Scenario and Cause Study

- **processes operate concurrently** (cooperation and competition)
  - increased utilization of shared resources (e.g., CPU)
- mandatory: **synchronisation** primitives are used to coordinate processes
  - data races induce the need of synchronisation methods
- concept: **passive waiting** by blocking and waking processes
  - resource usage and availability decides on process states
  - eliminates **busy waiting**, which is considered an anti-pattern (too expensive, error prone)
- solution with OS support: **semaphores** → implement **passive waiting** and enable:
  - unilateral synchronisation (→ e.g., producer/consumer)
  - multilateral synchronisation
    - interlaced producer/consumer, "prosumer"
    - mutual exclusion (e.g., exclusive access to a critical section)
- **waiting mechanisms** can lead to **deadlock** situations...




# Deadlocks (dt. Verklemmungen)







# Deadlocks Cause Study - Example



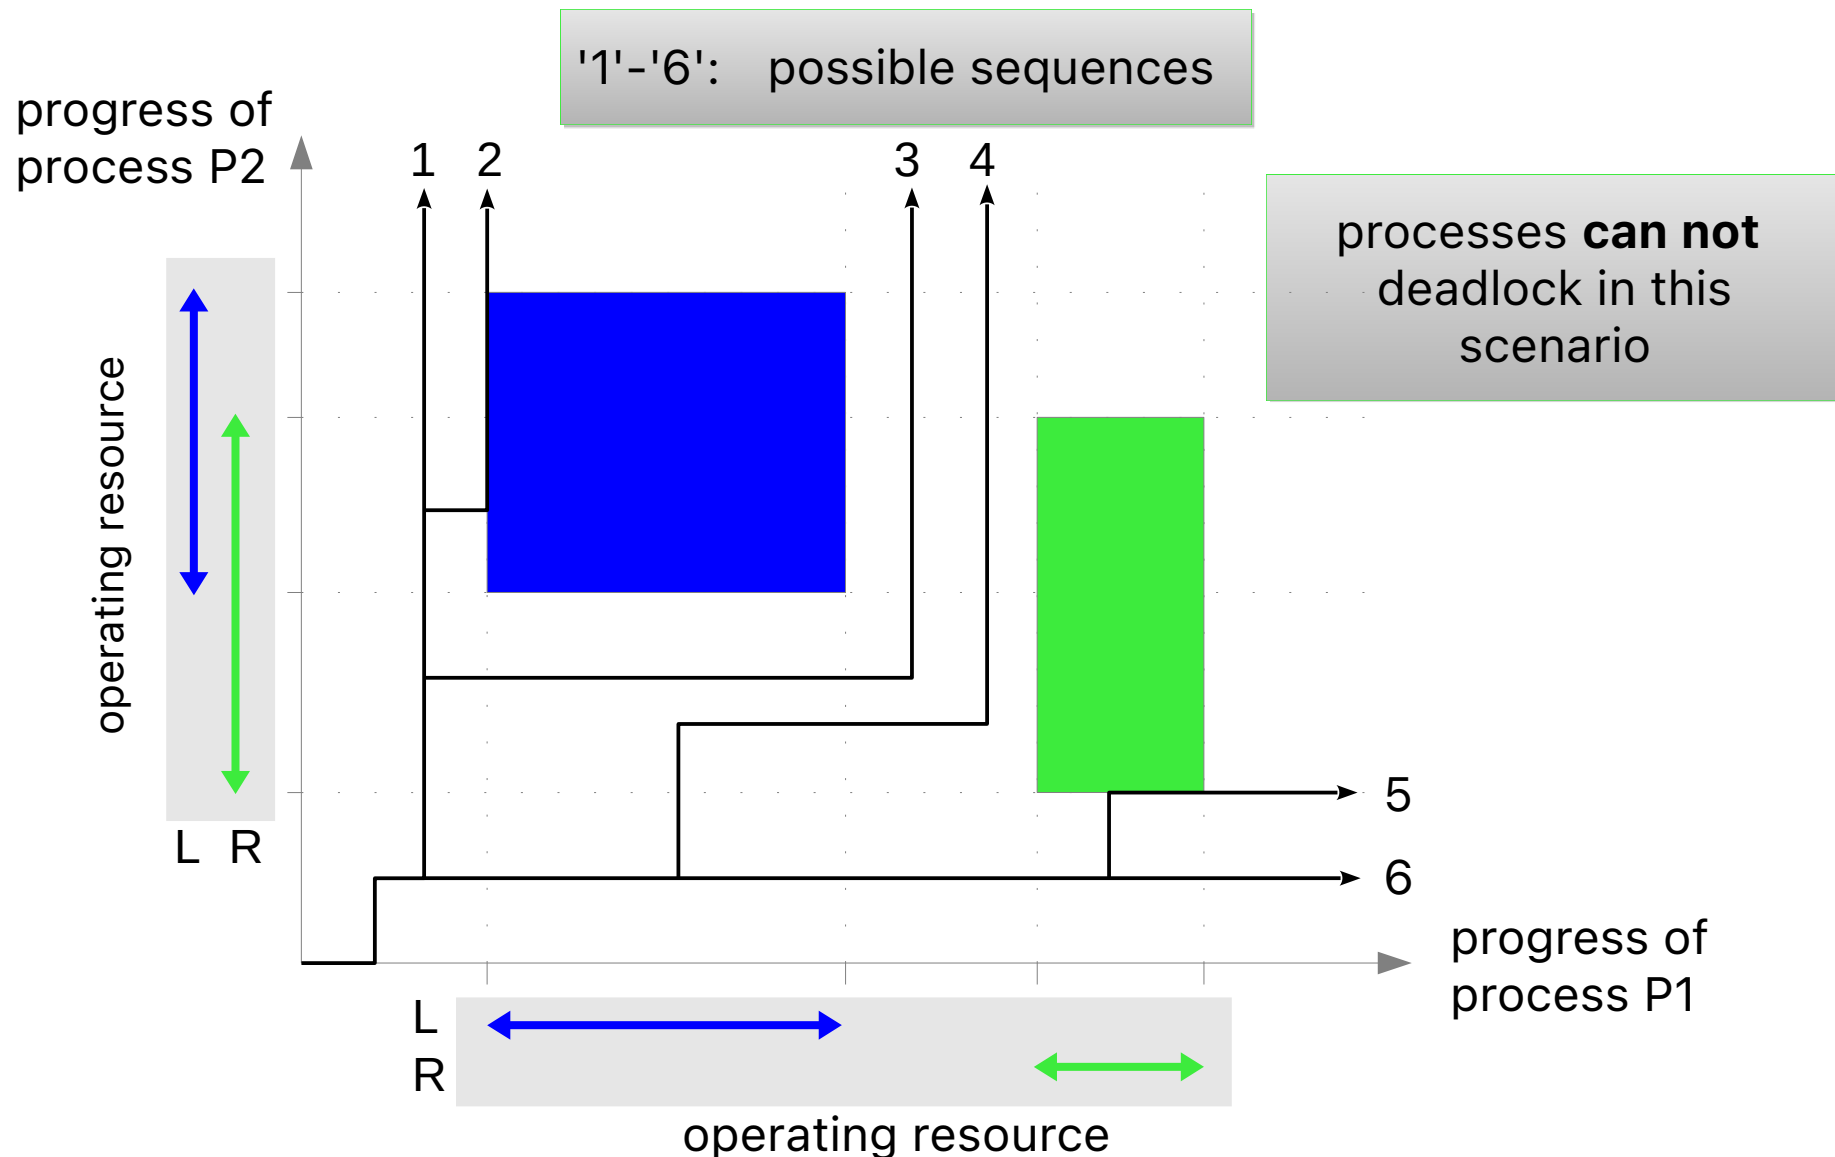
car 1 occupies L and requires R  
car 2 occupies R and requires L



# Deadlocks Cause Study - Abstract



# Deadlocks Cause Study - Abstract



# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Problem Scenario and Cause Study
- ▶ **Process Deadlocks**
  - ▶ Definition, Variants, and Conditions
  - ▶ Consumable and Non-consumable Resources
  - ▶ Resource Allocation Graphs
- ▶ Classic Deadlock Situation
  - ▶ The Dining Philosophers Problem
- ▶ Deadlock Counter Measures
  - ▶ Prevention, Avoidance, Detection and Recovery
- Summary and Outlook



# Deadlocks of Processes: Definition

- Deadly Embrace



Photo: David Maitland, National Geographic

- Definition (Tanenbaum and Woodhull):

*„A **set of processes is deadlocked** if **each process in the set is waiting** for an event that only **another process in the set can cause.**“*

# Deadlocks of Processes: Variants

- **deadlock** (first variant)
  - passive waiting
  - process state: **BLOCKED**
- **livelock** (second variant)
  - busy waiting (or „*lazy*“ *busy waiting*)
  - arbitrary process state (including **RUNNING**), but: no progress!
- by comparison, deadlocks are the lesser evil
  - condition can be identified → prerequisite for „resolution“
  - extremely high system load due to busy waiting

# Deadlocks of Processes: Conditions

For deadlocks to occur, **all** of the following conditions must be met ("necessary conditions"):

- 1. mutual exclusion condition:** exclusive allocation of resources
  - only one process may use a resource at a time
- 2. hold-and-wait condition:** follow-up request for resources
  - processes currently holding resources can request new resources
- 3. no preemption condition:** no withdrawal of resources
  - granted resources must be explicitly released by the process holding them
- 4. circular wait condition:** processes wait for each other circularly
  - several processes, each of which is waiting for a resource held by the next member of the chain



# Deadlocks of Processes: Conditions

For deadlocks to occur, **all** of the following conditions must be met ("necessary conditions"):

- 1. mutual exclusion condition:** exclusive allocation of resources
    - only one process may use a resource at a time
  - 2. hold-and-wait condition:** follow-up request for resources
    - processes currently holding resources can request new resources
  - 3. no preemption condition:** no withdrawal of resources
    - granted resources must be explicitly released by the process holding them
- A deadlock **only occurs** when a **fourth condition** holds true at runtime:
- 4. circular wait condition:** processes wait for each other circularly
    - several processes, each of which is waiting for a resource held by the next member of the chain

# Operating Resources (dt. Betriebsmittel)

- operating resources are managed by the operating system and are made accessible to the processes
- types of operating resources:
  - **consumable operating resources**
    - are generated (produced) and destroyed (consumed) at runtime
    - examples: interrupt requests, signals, messages, data from input devices
    - typical access synchronisation: **unilateral synchronisation**
  - **reusable operating resources**
    - are allocated by processes for a certain time and then released
    - examples: CPU, main and background memory, I/O devices, system data structures such as files, process table entries
    - typical access synchronisation: **multilateral synchronisation, mutual exclusion**
- both resource types are prone to deadlocks

# Operating Resources - Reusable

- a deadlock occurs when two processes use a **reusable operating resource** that is requested by the other process
- example: a computer system has 200 GiB of main memory, two processes allocate the memory step by step (memory allocation is done block by block)

## Process 1

```
...  
allocate 80 GiB;  
...  
allocate 60 GiB;
```

## Process 2

```
...  
allocate 70 GiB;  
...  
allocate 80 GiB;
```


If both processes issue their first request before memory is requested, a deadlock is unavoidable.

# Operating Resources - Consumable

- a deadlock occurs when two processes are waiting for a **consumable operating resource** to be produced by the other
- example: synchronisation signals are "sent" between two processes using the **wait** and **signal** semaphore operations

## Process P1

```
semaphore s1 = {0, NULL};  
wait (&s1);  
... // consume (from P2)  
... // produce (for P2)  
signal (&s2);
```



## Process P2


```
semaphore s2 = {0, NULL};  
wait (&s2);  
... // consume (from P1)  
... // produce (for P1)  
signal (&s1);
```

Each process waits for a synchronisation signal from the other, but this cannot be sent because it is blocked itself.


# Resource Allocation Graphs (RAGs)

(dt. **Betriebsmittelbelegungsgraphen**)

- to visualise and automatically detect deadlock situations Resource Allocation Graphs (RAGs) are used
- RAGs describe a current system state
  - **nodes:** processes and resources
  - **edges:** indicate an occupation or a request





operating resource  $R_1$  is **requested**  
by process  $P_1$



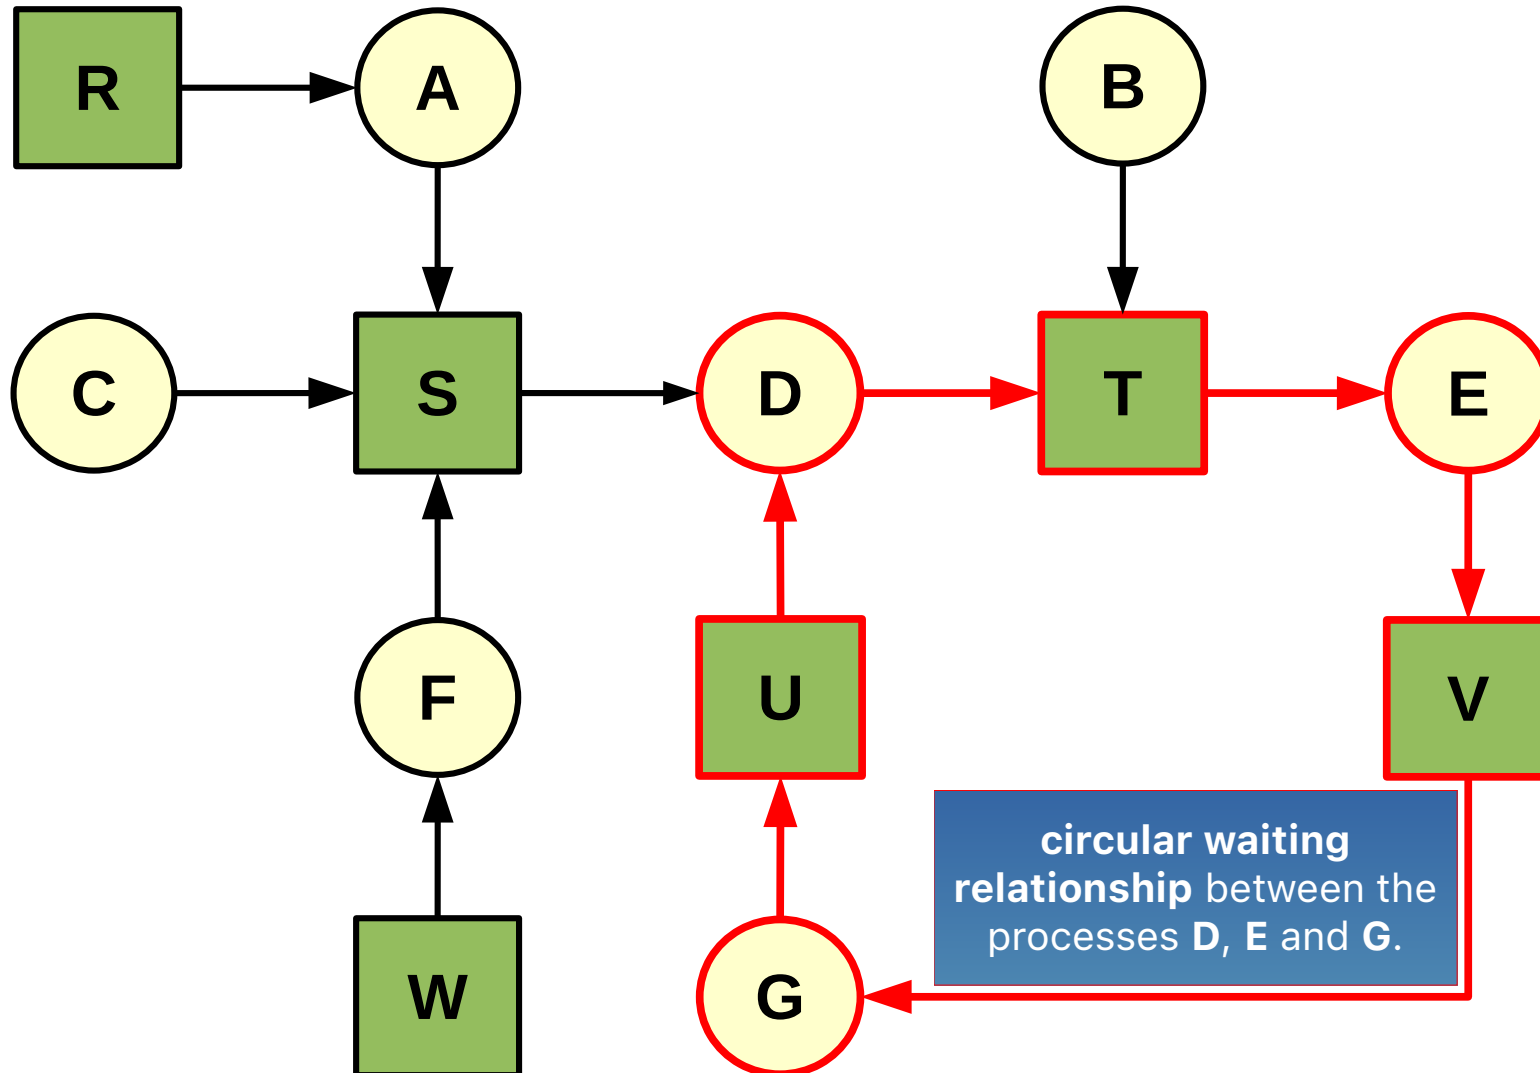
process  $P_2$  **has occupied** the  
operating resource  $R_2$

# Resource Allocation Graphs (RAGs)

- situation: there are 7 processes (**A** to **G**) and 6 resources (**R** to **W**)
  - is there circular waiting?
  - which processes are involved?
- **current state:**
  - **A** occupies **R** and requests **S**
  - **B** does not occupy anything, but requests **T**
  - **C** does not occupy anything, but requests **S**
  - **D** occupies **U** and **S** and requests **T**
  - **E** occupies **T** and requests **V**
  - **F** occupies **W** and requests **S**
  - **G** occupies **V** and requests **U**



# Resource Allocation Graphs (RAGs)



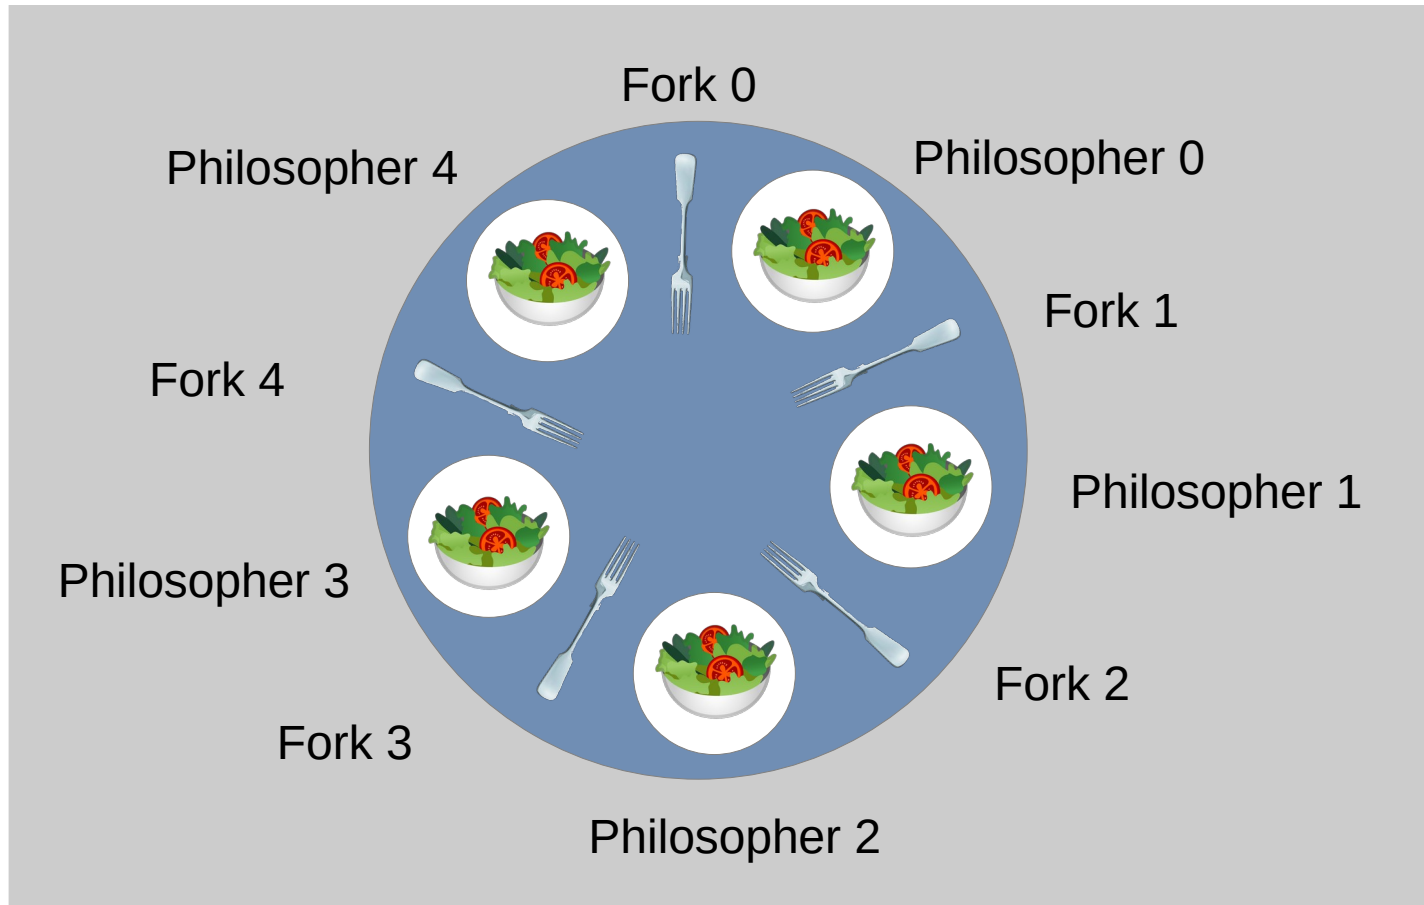


# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Problem Scenario and Cause St
- ▶ Process Deadlocks
  - ▶ Definition, Variants, and Conditions
  - ▶ Consumable and Non-consumable Resources
  - ▶ Resource Allocation Graphs
- ▶ Classic Deadlock Situation
  - ▶ The Dining Philosophers Problem
- ▶ Deadlock Counter Measures
  - ▶ Prevention, Avoidance, Detection and Recovery
- Summary and Outlook



# The Dining Philosophers Problem



The Dining Philosophers Problem (by E. W. Dijkstra)

Five philosophers, who have nothing else to do but think and eat, sit at a round table.

Thinking makes you hungry - so every philosopher will also eat. For this, however, a philosopher always needs **both** forks lying next to his plate.

**Philosopher**  
**Fork**

→ **Process**

→ **Operating Resource**

# Deadlocked Philosophers?

The first three necessary conditions are met:

- ✓ ■ **mutual exclusion**
  - for reasons, the philosophers are not allowed to share forks
- ✓ ■ **hold and wait**
  - the philosophers are so preoccupied with their thoughts before eating that they can neither really grasp their forks at the same time, nor do they think of putting a fork down again
- ✓ ■ **no preemption**
  - snatching the fork from another philosopher is out of the question

? **But:** does a deadlock really happen?

# Dining Philosophers - Attempt 1

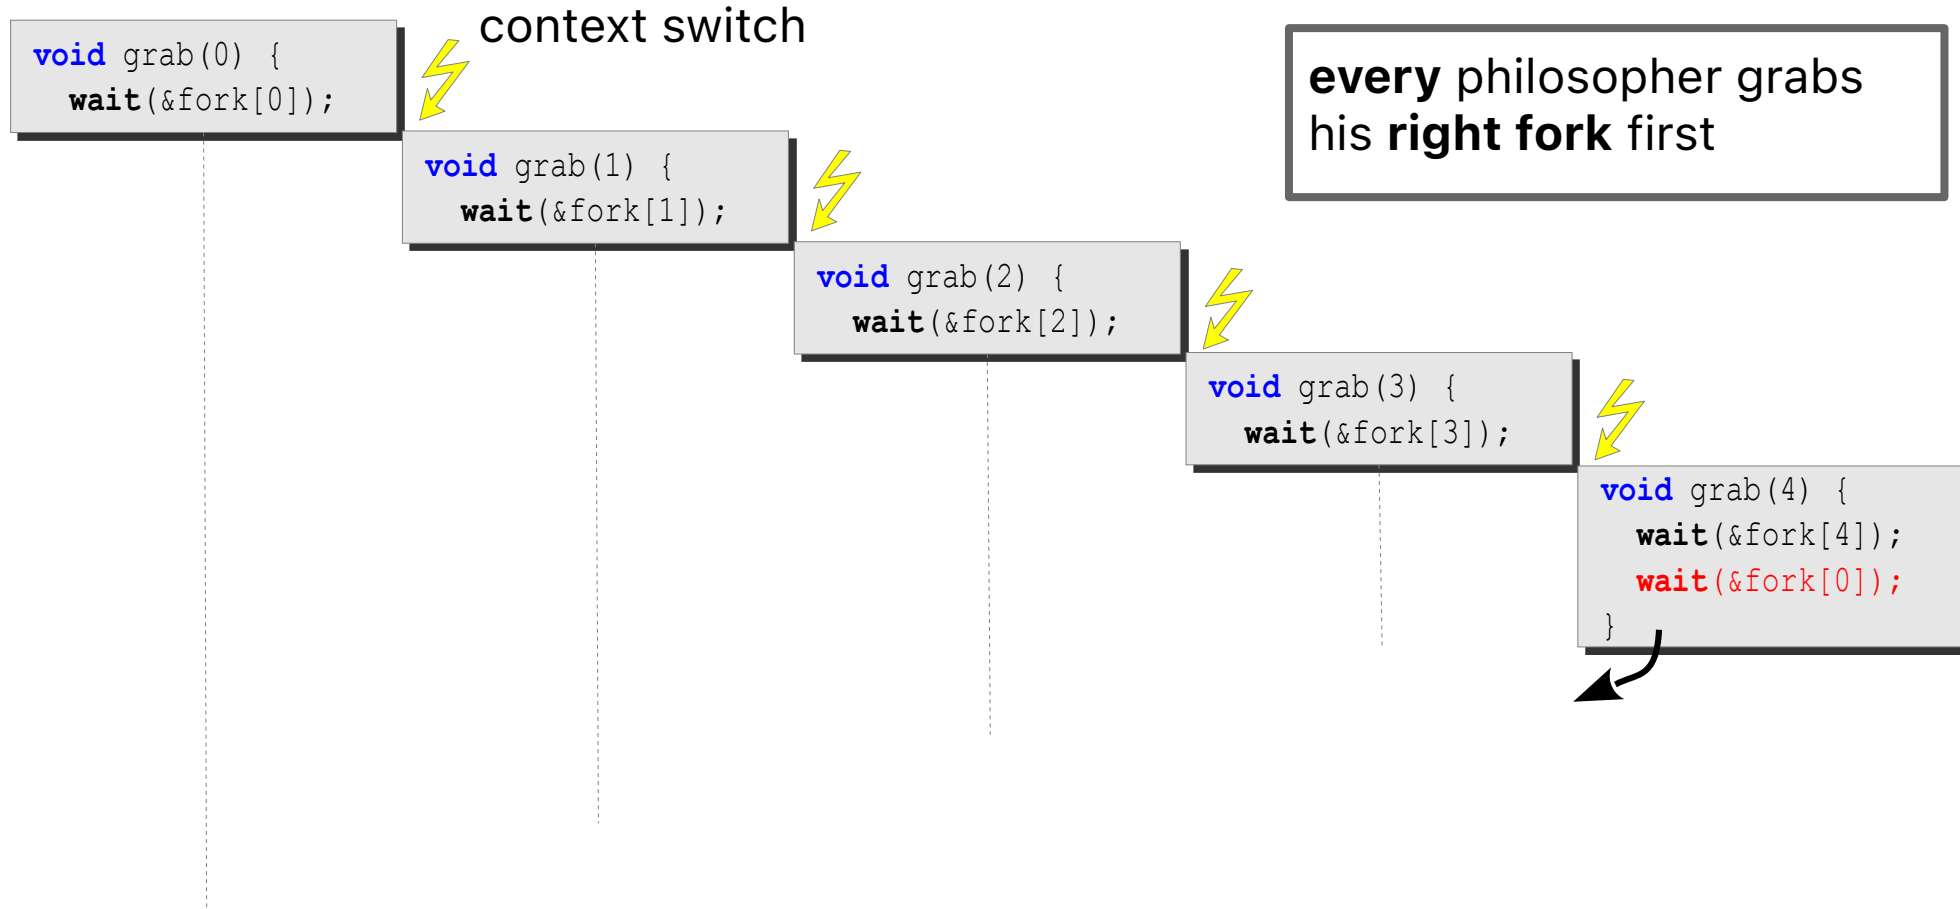
```
void phil (int who) {  
    while (1) {  
        think();  
        grab(who);  
        eat();  
        drop(who);  
    }  
}
```

```
void think () { ... }  
void eat    () { ... }
```

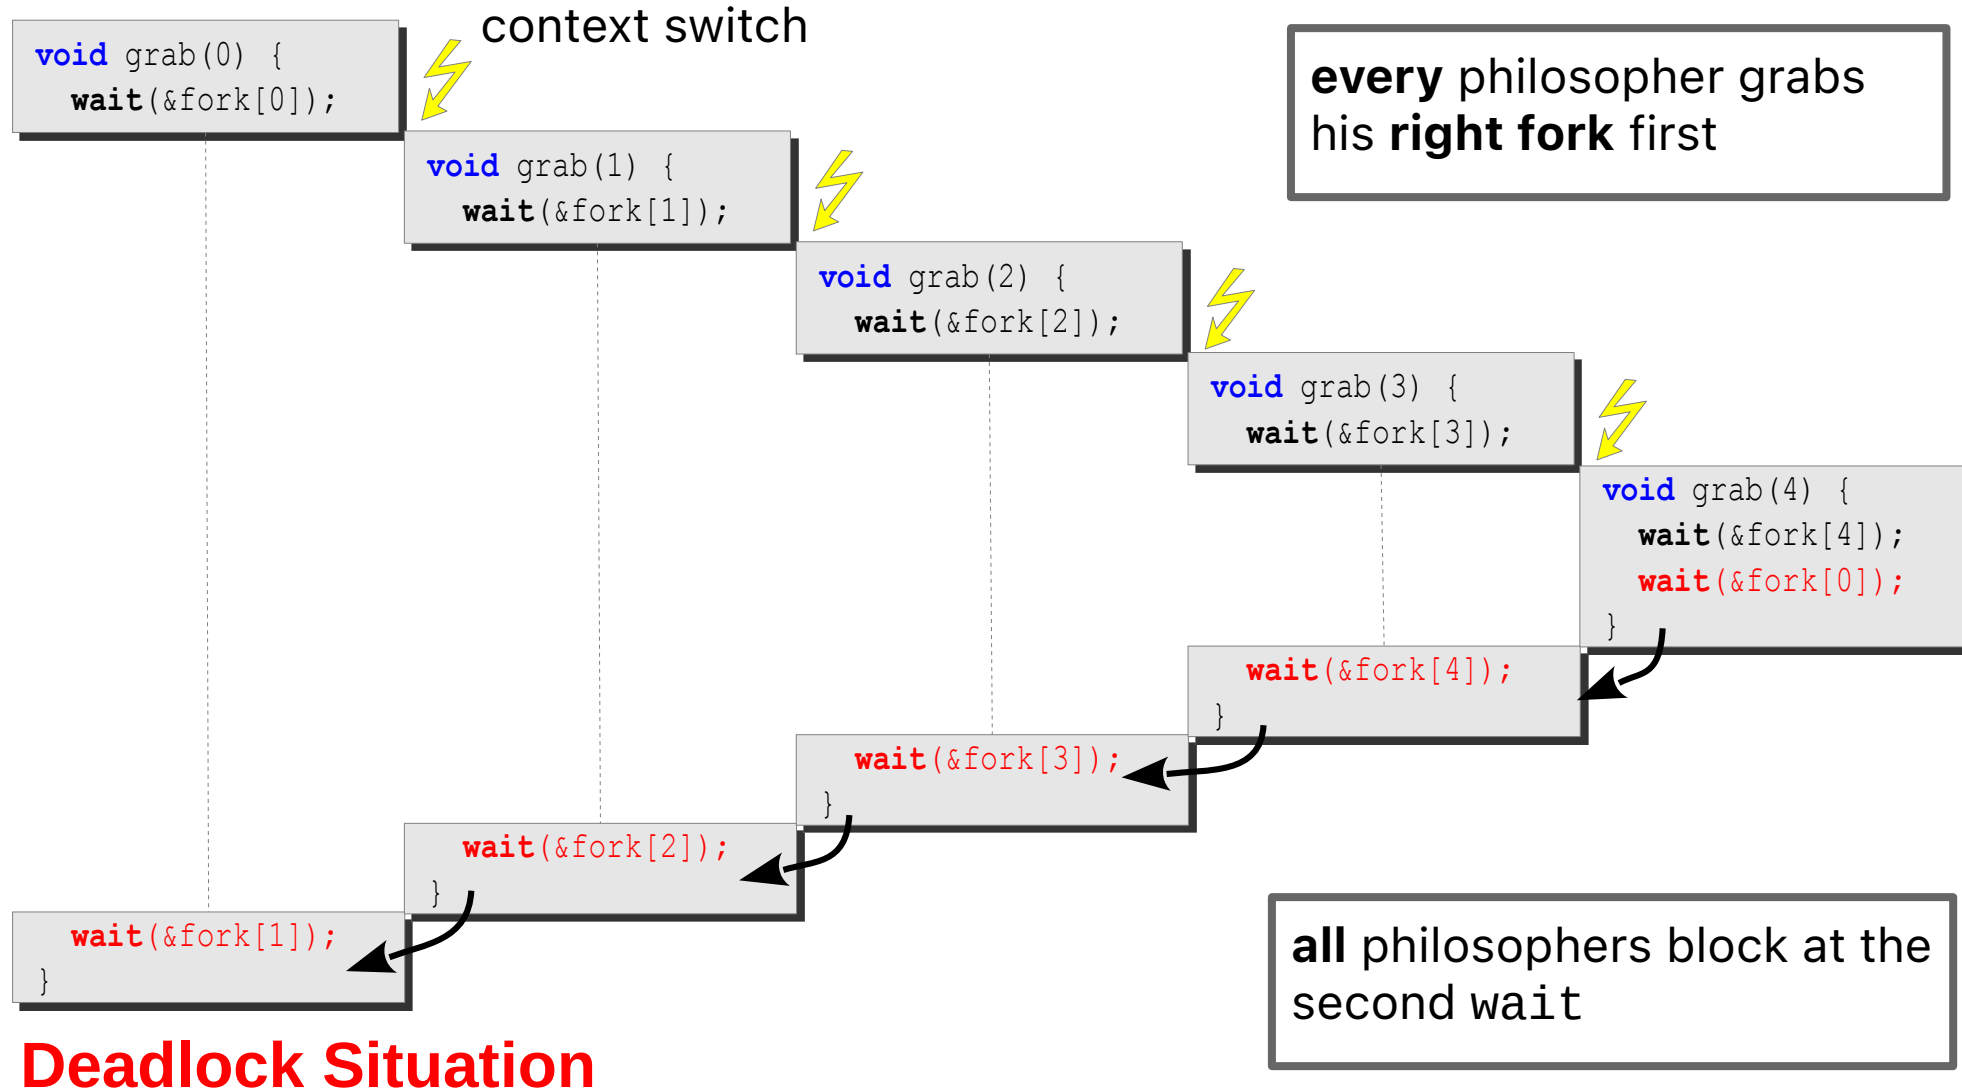
```
semaphore fork[NPHIL] = {  
    {1, NULL}, ...  
};  
  
void grab (int who) {  
    wait(&fork[who]);  
    wait(&fork[(who+1)%NPHIL]);  
}  
  
void drop (int who) {  
    signal(&fork[who]);  
    signal(&fork[(who+1)%NPHIL]);  
}
```

With the help of a semaphore, **mutual exclusion** is guaranteed when accessing the forks. Each philosopher first takes his right fork and then his left fork.

# Attempt 1: Prone to Deadlock Situations



# Attempt 1: Prone to Deadlock Situations




Att

```
void gr  
wait (
```

```
wait (
```

Dead

## Resource Allocation Graph



```
) ;  
) ;
```

he



# Dining Philosophers - Attempt 2

```
semaphore mutex = {1, NULL};  
  
void grab (int who) {  
    wait(&mutex);  
    wait(&fork[who]);  
    wait(&fork[(who+1) % NPHIL]);  
    signal(&mutex);  
}
```

The problem of the first implementation were context switches between the 1st and 2nd wait, in fact, a **critical section**.

The second implementation protects this critical section by mutual exclusion.

- Is it deadlock free?
- Is the second attempt a solid solution?

# Dining Philosophers - Attempt 2

```
semaphore mutex = {1, NULL};

void grab (int who) {
    wait(&mutex);
    wait(&fork[who]);
    wait(&fork[(who+1) % NPHIL]);
    signal(&mutex);
}
```

The problem of the first implementation were context switches between the 1st and 2nd wait, in fact, a **critical section**.

The second implementation protects this critical section by mutual exclusion.



■ Is it deadlock free? **Yes**

- max. 1 process can wait for a fork (cycle needs 2!)
- a process waiting for **mutex** has no fork



■ Is the second attempt a solid solution? - **No!**

- when  $\text{philo}_{\text{who}}$  eats,  $\text{philo}_{\text{who}+1}$  blocks the entry of the critical section and **all other** philosophers **block**, too.
- **result: low degree of concurrency** and **inefficient use of available resources**

# Dining Philosophers - Attempt 3

```
const int N = 5; /* number of dining philosophers */
semaphore mutex = {1, NULL}; /* mutual exclusion */
semaphore s[N] = {{0, NULL},...}; /* one semaphore for each philosopher */
enum {THINKING, EATING, HUNGRY } state[N]; /* philosopher states */

int left(i) { return (i+N-1)%N; } /* index left neighbor */
int right(i) { return (i+1)%N; } /* index right neighbor */
```

```
void test (int i) {
    if (state[i] == HUNGRY && state[left(i)] != EATING &&
        state[right(i)] != EATING) {

        state[i] = EATING;
        signal(&s[i]);
    }
}
```

```
void grab(int i) {
    wait(&mutex);
    state[i] = HUNGRY;
    test(i);
    signal(&mutex);
    wait(&s[i]);
}
```


```
void drop(int i) {
    wait(&mutex);
    state[i] = THINKING;
    test(left(i));
    test(right(i));
    signal(&mutex);
}
```

This solution is  
deadlock-free and has a  
maximum degree of  
concurrency

# Dining Philosophers - Discussion

- **for specific problems:** there are usually many different ways to ensure deadlock freedom
  - solutions differ in the degree of possible concurrency
  - if a solution is too restrictive, operating resources will temporarily lie idle unnecessarily
- **in general:** representative example of deadlock problems in the management of indivisible resources
  - based on E.W. Dijkstra (1965)
  - established standard scenario for evaluating and illustrating operating system and language mechanisms for concurrent programming

# Agenda

- ▶ Recap
  - ▶ Organizational Matters
  - ▶ Problem Scenario and Cause Study
  - ▶ Process Deadlocks
    - ▶ Definition, Variants, and Conditions
    - ▶ Consumable and Non-consumable Resources
    - ▶ Resource Allocation Graphs
  - ▶ Classic Deadlock Situation
    - ▶ The Dining Philosophers Problem
  - ▶ Deadlock Counter Measures
    - ▶ Prevention, Avoidance, Detection and Recovery
  - Summary and Outlook
- 



# Deadlock Prevention

(dt. **Verklemmungsvorbeugung**)

- **indirect methods** invalidate (at least) one of the conditions 1-3
    1. use of non-blocking process synchronisation methods
    2. design resource requests in an indivisible (atomic) way
    3. withdrawal of operating resources through **virtualisation**
      - virtual memory, virtual devices, virtual processor cores
  - **direct methods** invalidate condition 4
    4. linear/total order of operating resources:
      - operating resource  $R_i$  can only be allocated before resource  $R_j$ , if  $i$  is in a linear order directly before  $j$  (i.e.,  $i < j$ ).
- indirect/direct methods: rules, that prevent the occurrence of deadlocks
- methods that take effect at **design/implementation time**

# Deadlock Avoidance

(dt. **Verklemmungsvermeidung**)

- preventing circular waiting (**at run-time**) through strategic measures:
  - none of the first three necessary conditions is invalidated
  - ongoing **resource requirement analysis** eliminates circular waiting
- requests for operating resources of the processes are to be controlled:
  - „**safe state**“ must always be maintained:
    - there is a process sequence in which each process can cover its maximum resource requirement
  - „**unsafe states**“ must be **avoided**:
    - allocation rejection in the event of uncovered operating resource requirements
    - do not serve requesting processes or suspend them at an early stage
- problem: a priori knowledge of maximum operating resource requirements of processes is necessary



# Deadlock Avoidance: Safe and Unsafe States

(using

■ i

■ c

F

2

## „Banker's Algorithm“

- management of process/operating resource matrices for **current occupancy** and **maximum occupancy**
- function for **finding a process sequence** in which the resources do not run out even if the "credit limit" is exhausted completely
- anticipatory application of this function in the event of fulfilling resource allocation requests

(c.f., Tanenbaum)

ents

nents

ch,

# Deadlock Avoidance: Safe and Unsafe States

(using the example of multiple existing resources)

- **initial state:** a UNIX system with **12 shared memory segments**
  - process P1 requires max. 10 segments, P2: 4 segments, P3: 9 segments
- **current state:**

P1 occupies 6 segments, P2 and P3 occupy 2 segments each,  
2 segments are still available (free)

  - P3 requests 1 segment, 1 would remain free → **unsafe state**
    - P3's request is rejected, P3 has to wait
  - P1 requests 2 segments, no free segment remains → **unsafe state**
    - the request of P1 is rejected, P1 has to wait
  - **safe process sequence:** P2 → P1 → P3

# Deadlock Detection

(dt. **Deadlockerkennung**)

- deadlocks are (silently) accepted („**ostrich algorithm**“) . . .
  - nothing in the system prevents the occurrence of circular wait conditions
  - *none* of the four conditions is invalidated
- approach: create **waiting graph** and search cycles →  **$O(n)$** 
  - too frequent verification wastes resources/computational power
  - too infrequent verification leaves operating resources lying idle
- **search for cycles** occurs mostly at large time intervals when:
  - serving operating resource requests takes too long
  - the utilisation of the CPU decreases despite process increase
  - the CPU has already been idle for a very long period of time

# Deadlock Recovery

recovery phase *after* the detection phase


- **terminate processes** and thus free up operating resources
  - stopping deadlocked processes step by step (huge efforts)
    - start with the "most effective victim" – who is it?
  - terminate **all** deadlocked processes (huge damage)
- **revoke operating resources** and begin with the „most effective victim“ – again, who is it?
  - reverse or restart the identified process
    - transactions, checkpointing/recovery (huge effort)
  - starvation of the processes that have been reversed must be avoided
  - also, beware of livelocks!
- walking a tightrope between harm and expense:
  - damage is unavoidable, potential for (severe) follow-up damage

# Discussion on Countermeasures

- deadlock *prevention, avoidance, detection, recovery*
- methods for avoidance/detection are less relevant in practice in the context of operating systems
  - methods are hard to implement, too costly and therefore not applicable
  - moreover, the prevalence of sequential programming makes these methods little necessary
- risk of deadlocks can be solved by virtualisation of operating equipment
  - processes use/occupy **logical operating resources** only
  - goal: withdraw **physical operating resources** from the processes (without their knowledge) at critical moments
  - with this, condition of preemption is defused
- deadlock prevention methods are more relevant

# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Problem Scenario and Cause Study
- ▶ Process Deadlocks
  - ▶ Definition, Variants, and Conditions
  - ▶ Consumable and Non-consumable Resources
  - ▶ Resource Allocation Graphs
- ▶ Classic Deadlock Situation
  - ▶ The Dining Philosophers Problem
- ▶ Deadlock Counter Measures
  - ▶ Prevention, Avoidance, Detection and Recovery
- Summary and Outlook



# ▶▶ Summary and Outlook

## ■ summary

- **process deadlocks:** mutual blocking of concurrent but independent processes
  - deadly embrace: deadlocks and livelocks
  - four deadlock conditions must hold true – invalidate (at least) a single condition to avoid deadlocks
  - occurs with reusable and consumable resources
- **deadlock handling**
  - dining philosophers and discussion of (potential) solutions
  - countermeasures: deadlock prevention, avoidance, detection and resolution

## ■ outlook: memory management

- main memory is a fundamental operating resource to processes
- memory allocation and memory management for multi-program operation, address mapping

# References and Acknowledgments

## Lecture

- ▶ Systemnahe Programmierung in C (SPiC), Betriebssysteme (Jürgen Kleinöder, Wolfgang Schröder-Preikschat)
- ▶ Betriebssysteme und Rechnernetze (Olaf Spinczyk, Embedded Software Systems Group, Universität Osnabrück)

## Teaching Books and Reference Book

- [1] Avi Silberschatz, Peter Baer Galvin, Greg Gagne: *Operating System Concepts*, John Wiley & Sons, 2018.
- [2] Andrew Tanenbaum, Herbert Bos: *Modern Operating Systems*, Pearson, 2015.
- [3] Wolfgang Schröder-Preikschat: *Grundlage von Betriebssystemen – Sachwortverzeichnis*, 2023.  
<https://www4.cs.fau.de/~wosch/glossar.pdf>