

# Operating Systems

Timo Hönig

Bochum Operating Systems and System Software (BOSS)

Ruhr University Bochum (RUB)

VII. Virtual Memory

May 24, 2023 (Summer Term 2023)



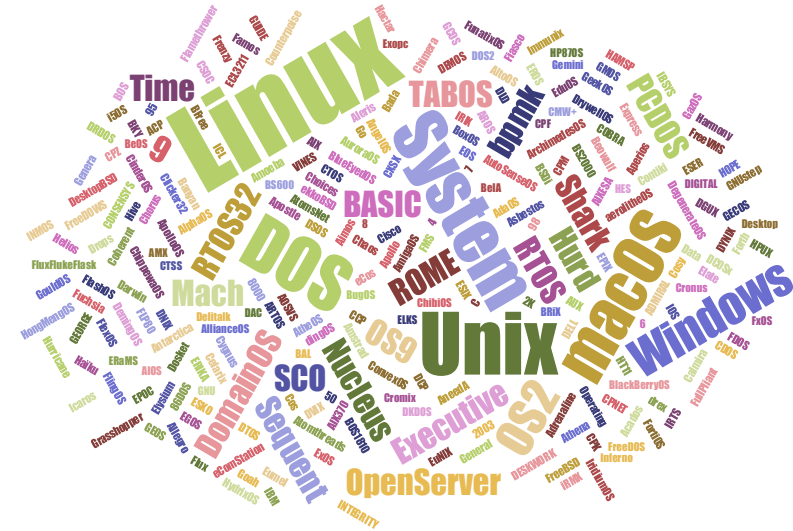
RUHR  
UNIVERSITÄT  
BOCHUM

RUB

[www.informatik.rub.de](http://www.informatik.rub.de)  
Chair of Operating Systems and System Software

# Agenda


- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ Demand Paging
- ▶ Page Replacement Algorithms
  - ▶ First In, First Out (FIFO),
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ Summary and Outlook



## Literature References

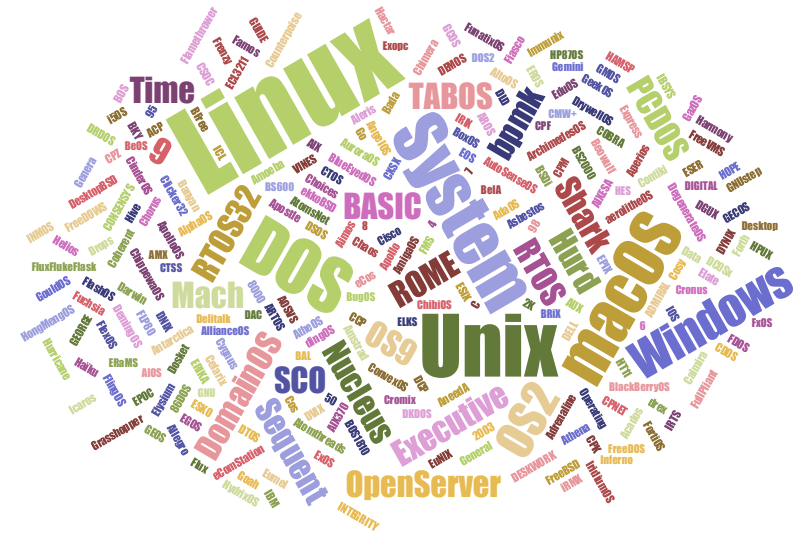
Silberschatz, Chapter 10

Tanenbaum, Chapter 3.3



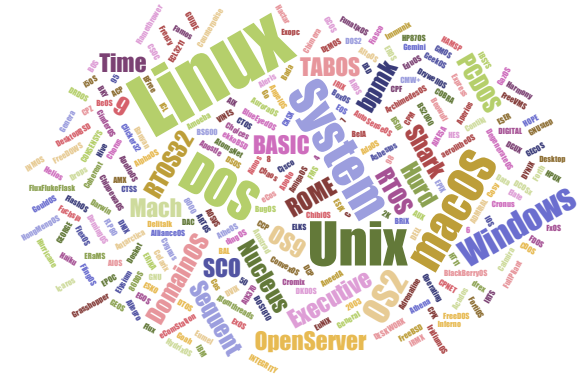
# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ Demand Paging
- ▶ Page Replacement Algorithms
  - ▶ First In, First Out (FIFO)
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ Summary and Outlook



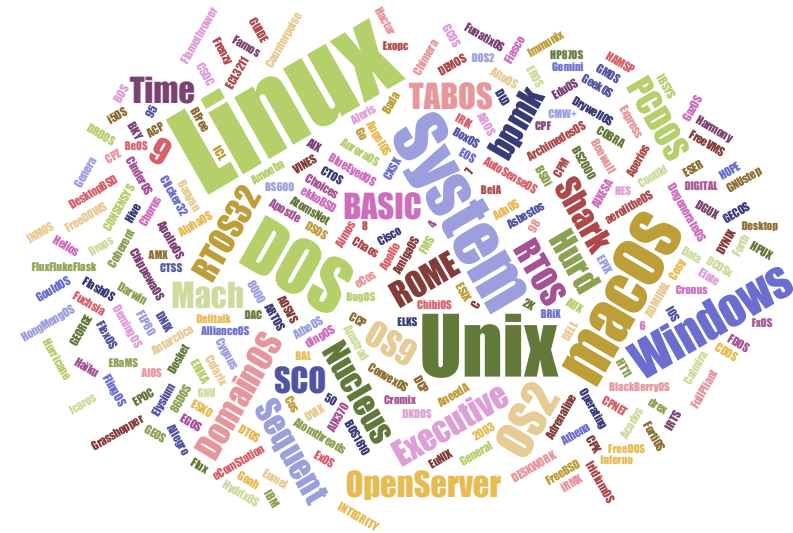
# Organizational Matters

- lecture
  - Wednesday, 10:15 – 11:45
  - format: synchronous, **hybrid**
    - in presence (Room H1D, Building ID)
    - online lecture (Zoom)
  - **exam:** August 7, 2023 (first appointment)  
September 25, 2023 (retest appointment)
- exercises: **group allocation almost complete**
  - make use of group work – for your own benefit!
- manage course material, asynchronous communication: Moodle
- <https://moodle.ruhr-uni-bochum.de/course/view.php?id=50698>



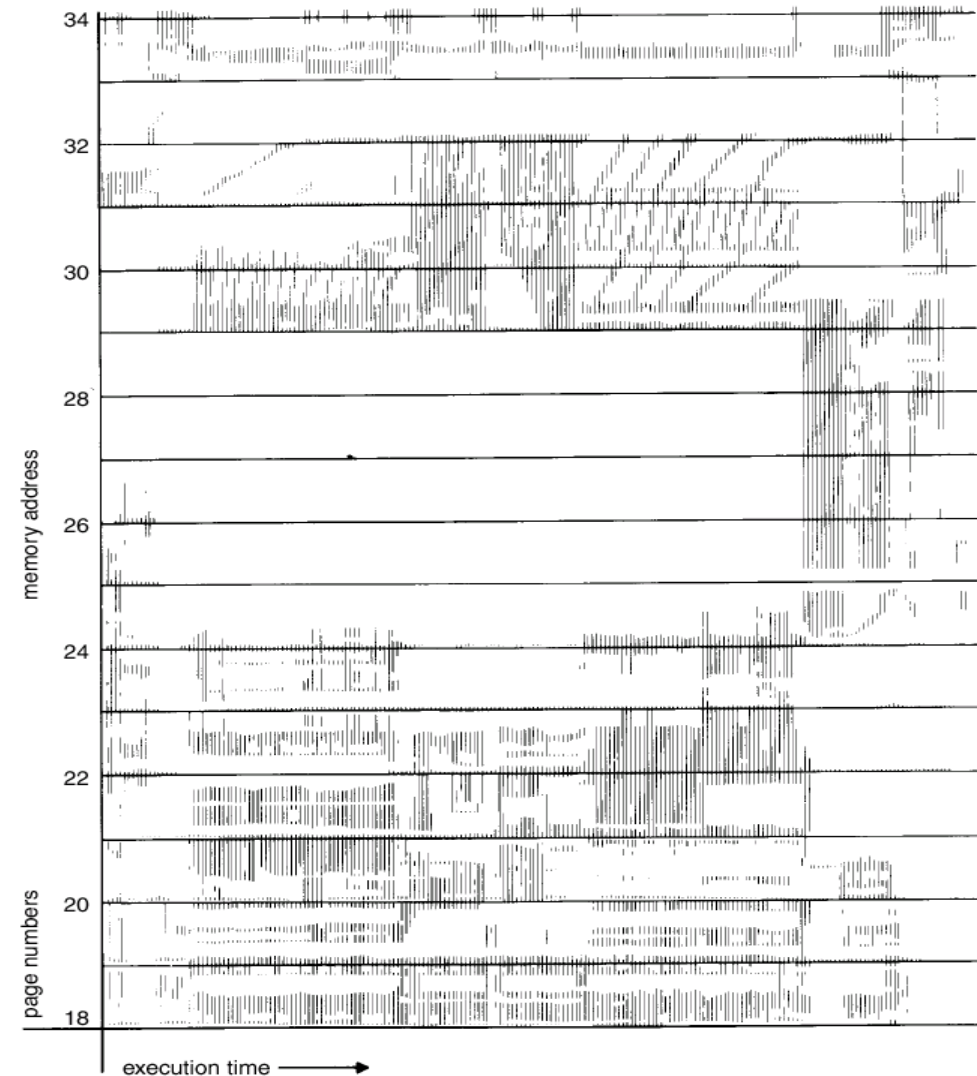
# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ Demand Paging
- ▶ Page Replacement Algorithms
  - ▶ First In, First Out (FIFO)
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ Summary and Outlook



# Locality of Reference – Memory Access

- single instructions require only a few memory pages
- even over longer periods of time: strong locality of reference
  - instructions are executed, for example, one after the other
- ➔ locality can be exploited if primary memory is not large enough to keep pages of all processes at the same time
  - for example, overlay technique



Source: Silberschatz, „Operating System Concepts“

# Virtual Memory – Main Idea

- add a system layer of indirection: decouple the memory requirement of processes from the system's main memory
  - processes do not need all memory locations equally often
    - certain functions are rarely or not at all used (e.g., error handling)
    - certain data structures are not fully populated
  - processes may require more memory than main memory is available
- main idea
  - delude the existence of a (unlimited) large main memory
  - swap out unused memory areas from primary to secondary storage
    - byte-addressable main memory (e.g., RAM) →
    - block-addressable background storage/memory (e.g., SSD, HDD)
  - provide the required memory areas on demand → demand paging

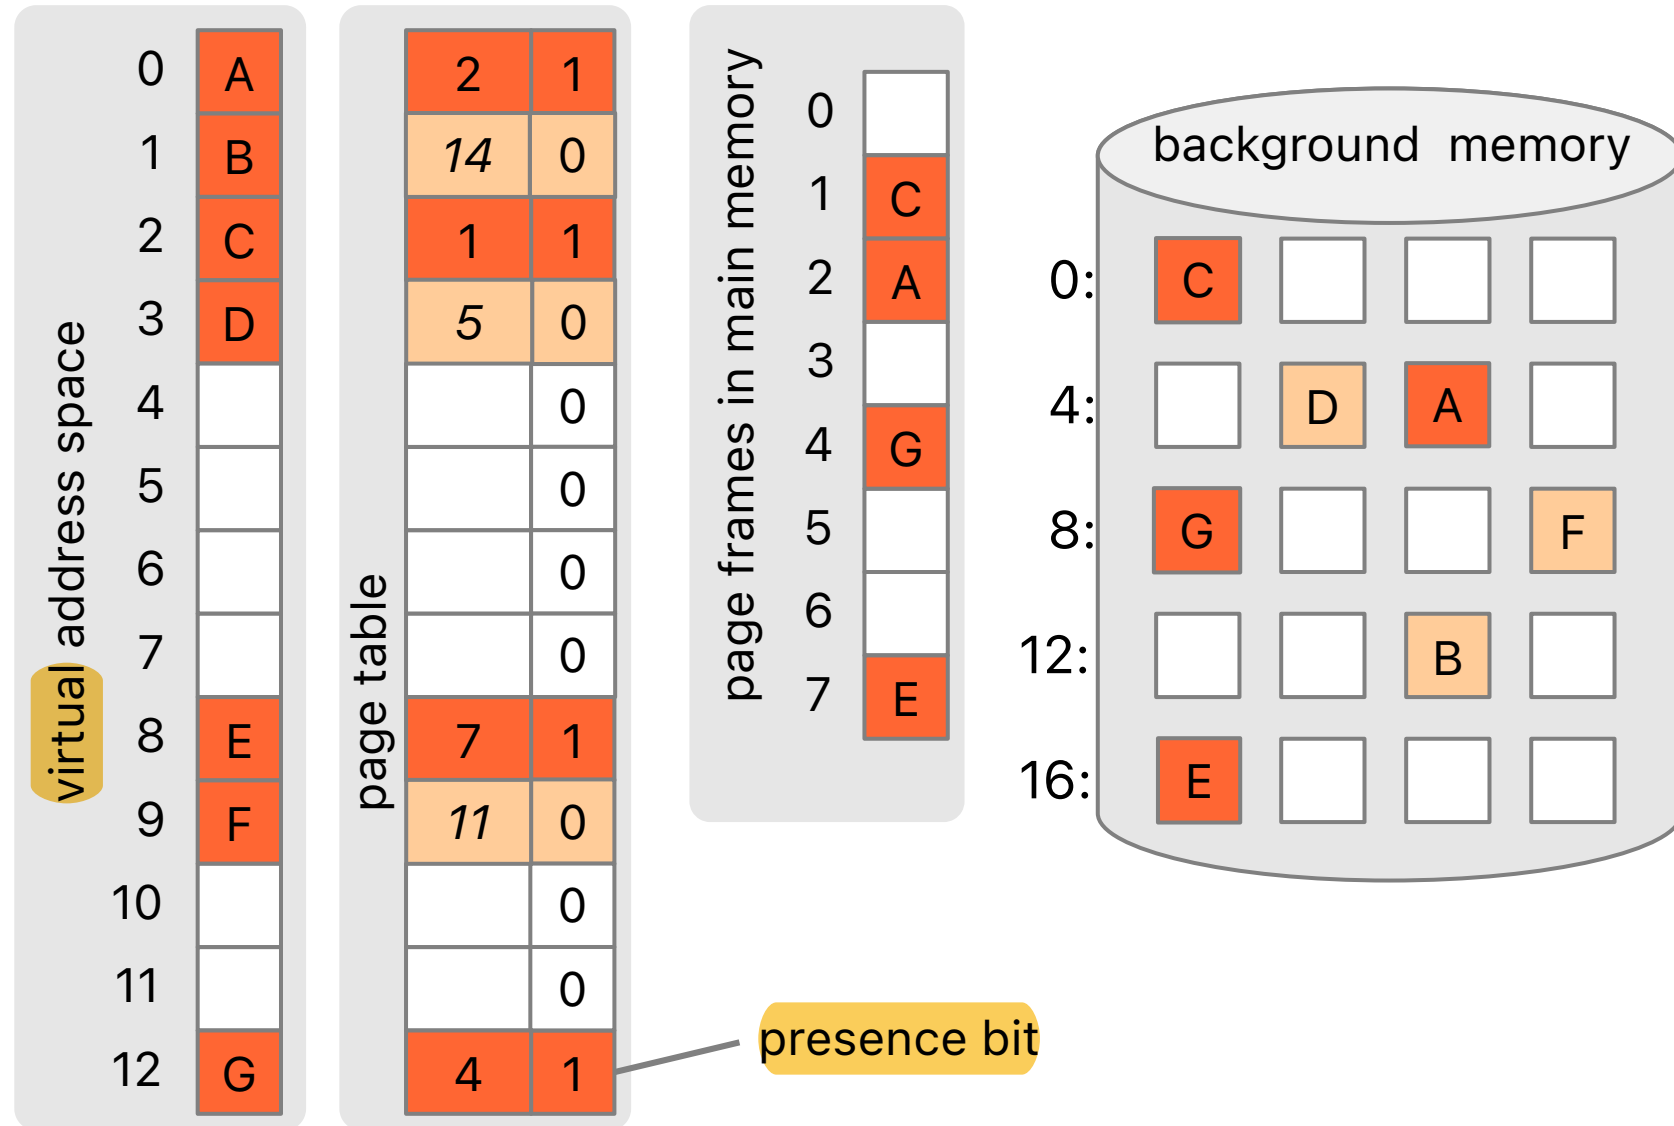


# Agenda

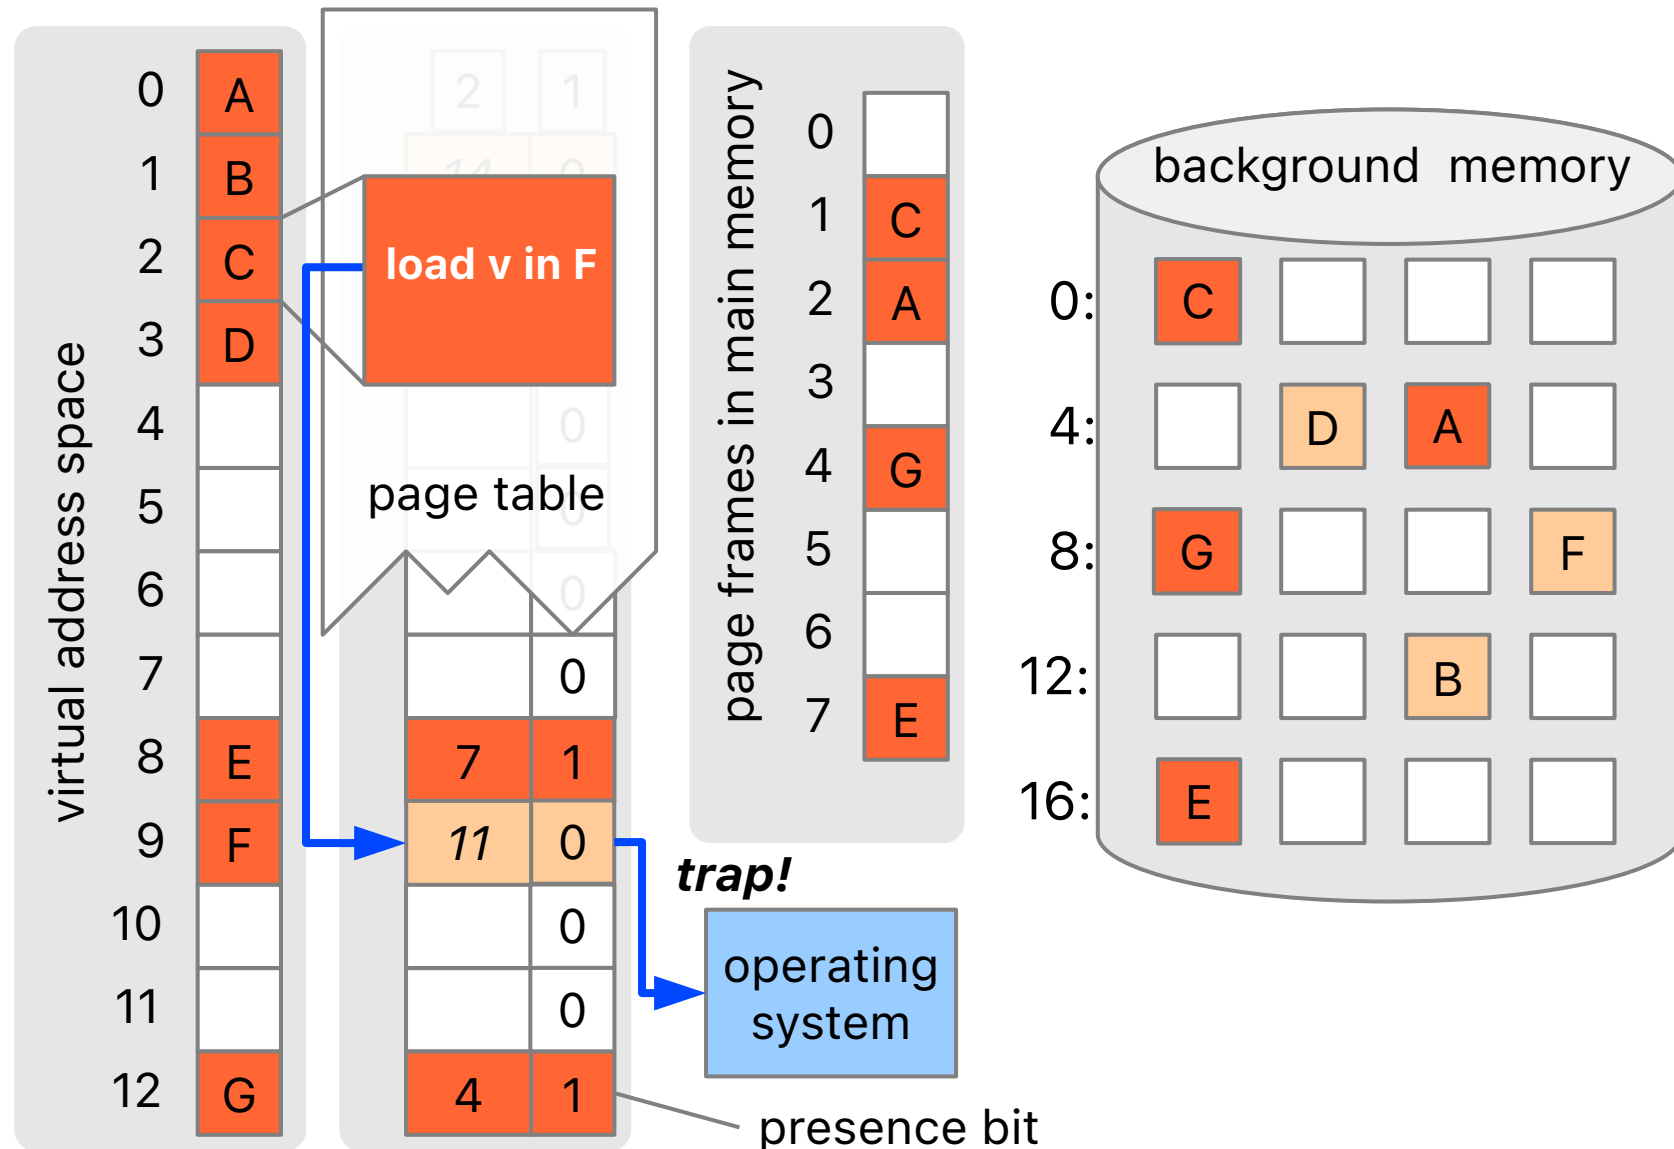
- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ **Demand Paging**
- ▶ Page Replacement Algorithms
  - ▶ First In, First Out (FIFO)
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ Summary and Outlook



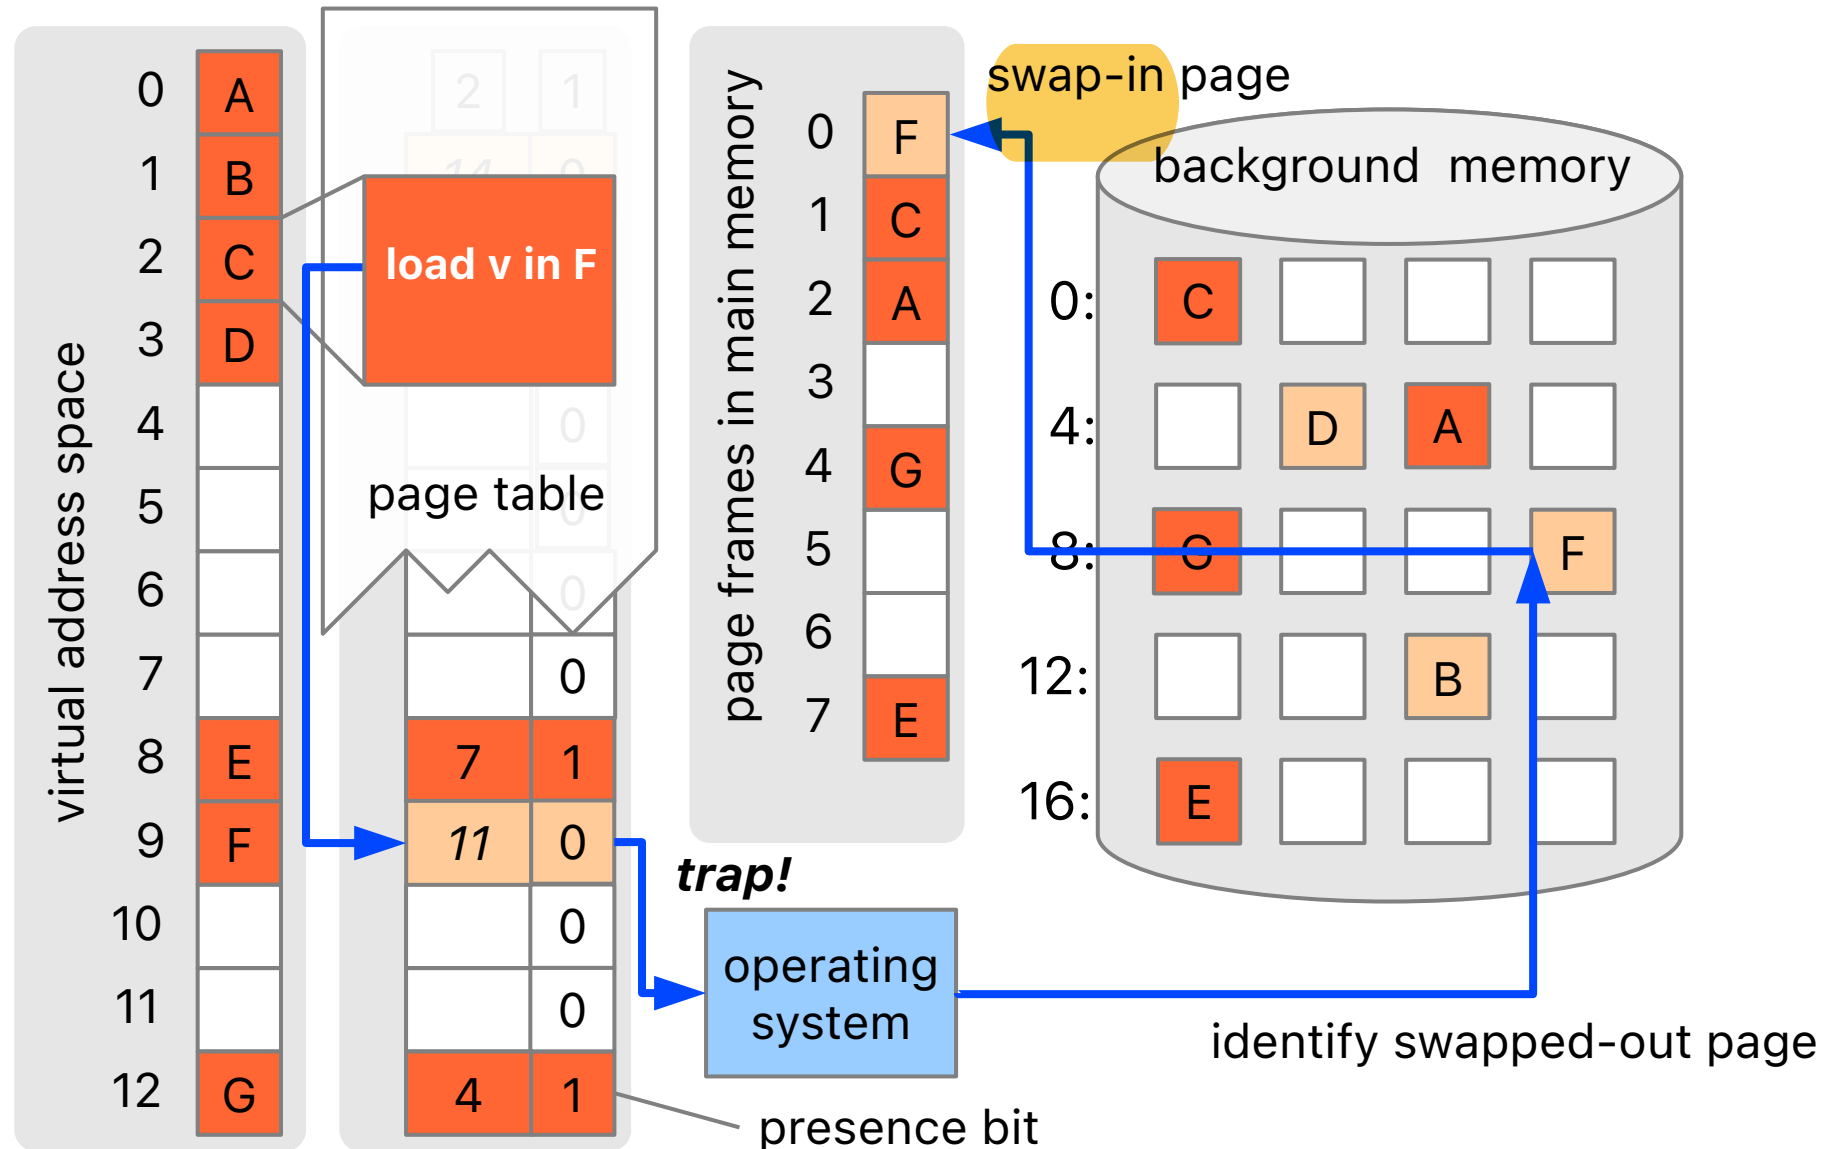
# Demand Paging – Provide Pages on Demand



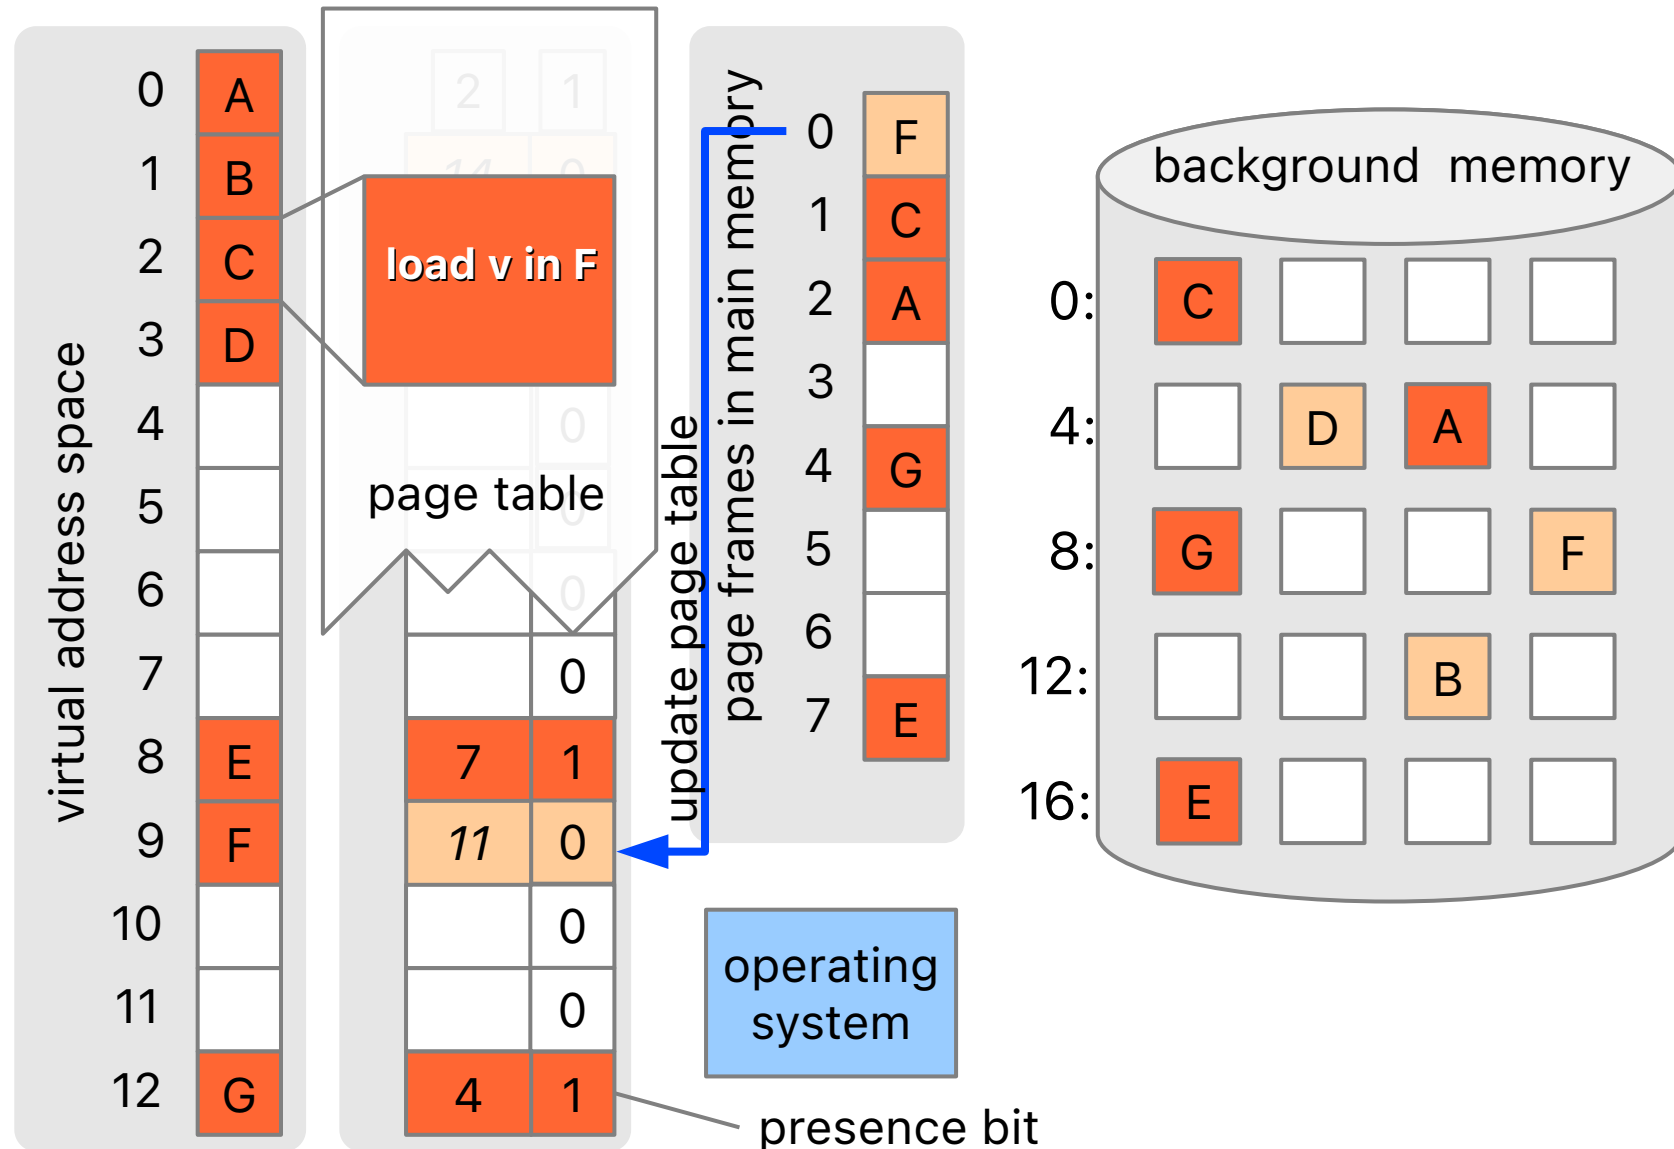
# Demand Paging – Page Fault Handling



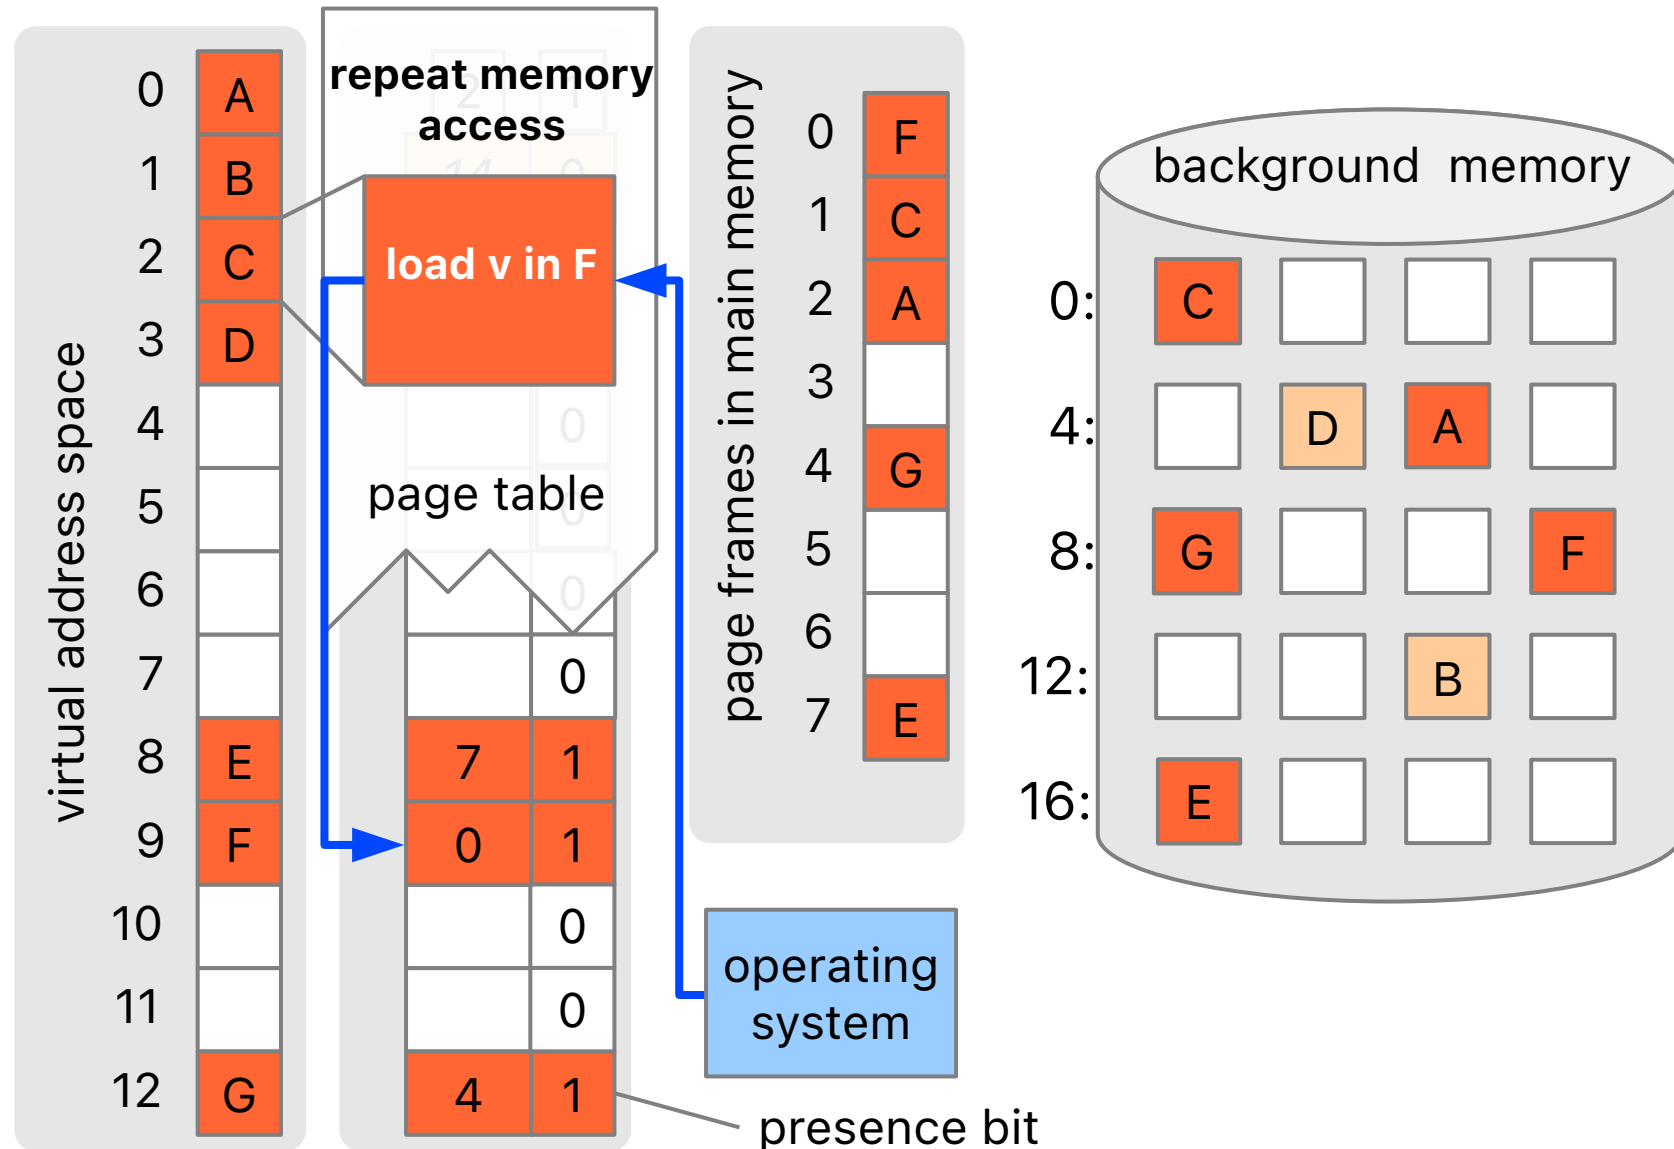
# Demand Paging – Page Fault Handling



# Demand Paging – Page Fault Handling



# Demand Paging – Page Fault Handling



# Demand Paging – Discussion

- operational costs of demand paging (i.e., performance impact)
    - without page fault
      - effective access time between 10 and 200 nanoseconds
    - with page fault
      - $p$ : probability of a page fault
      - first assumption: time to swap-in a page from background memory is 25 milliseconds (8 ms latency, 15 ms positioning time, 1 ms transmission time)
      - second assumption: main memory access time is 100 ns
      - effective access time:  
$$(1 - p) \cdot 100 + p \cdot 25000000 = 100 + 24999900 \cdot p$$
- rate of page faults must be *very low*
- $p$  close to zero
  - realistic only due to data locality

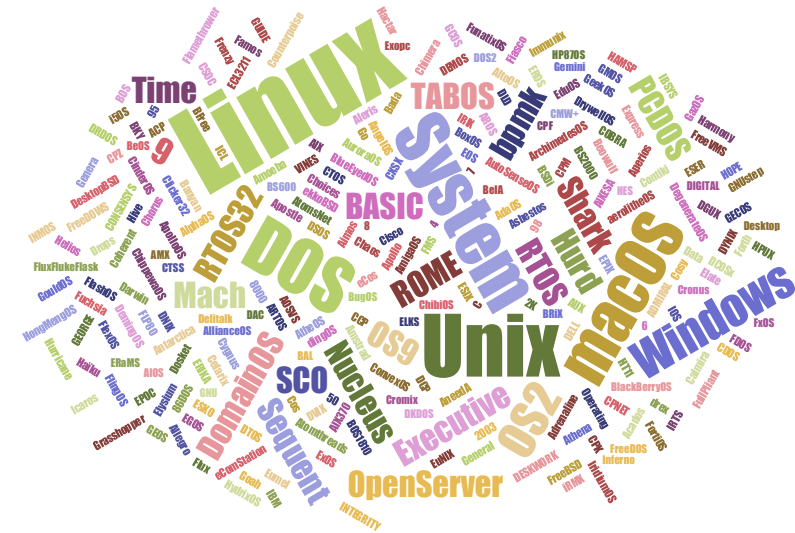
# Demand Paging – Characteristics

- operational concerns on process generation (e.g., fork)
  - **copy-on-write**
    - also easy to implement with paging MMU
    - finer granularity than with segmentation
  - program execution and memory accesses are interleaved
    - required pages are loaded only as they are needed
- page locking
  - lock pages that are used by DMA during trap (i.e., page fault handling)
  - necessary for input/output operations



# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ Demand Paging
- ▶ **Page Replacement Algorithms**
  - ▶ First In, First Out (FIFO)
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ Summary and Outlook



# Page Replacement Algorithms

- demand paging: need for free page frame(s)
- modus operandi:
  - page fault: exception handling via trap into the operating system
  - if no free page frame is available: swap-out a page
  - swap-in the required page
  - repeat memory access
- what to do if there is no free page frame?
  - a page must be swapped-out to the background memory to free up space for a new page in main memory
  - pages that have not been changed (dirty bit in page table)
  - if page was changed: process preemption requires swapping
- problem
  - which page ("victim") should be swapped-out?

# Page Replacement Algorithms

- consideration of replacement strategies and their effect on reference sequences
- reference sequence
  - sequence of page numbers that model the memory access behavior of a process
  - determination of reference sequences, for example, by recording the accessed addresses
    - reduce the recorded sequence to page numbers
    - group accesses to the same page in immediate succession
  - reference sequence: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	1											
	page frame 2												
	page frame 3												
control states (age per page frame)	page frame 1	0											
	page frame 2	>											
	page frame 3	>											

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	1	1										
	page frame 2		2										
	page frame 3												
control states (age per page frame)	page frame 1	0	1										
	page frame 2	>	0										
	page frame 3	>	>										

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1									
	page frame 2		<b>2</b>	2									
	page frame 3			<b>3</b>									
control states (age per page frame)	page frame 1	0	1	2									
	page frame 2	>	0	1									
	page frame 3	>	>	0									

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>								
	page frame 2		<b>2</b>	2	2								
	page frame 3			<b>3</b>	3								
control states (age per page frame)	page frame 1	0	1	2	0								
	page frame 2	>	0	1	2								
	page frame 3	>	>	0	1								

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4							
	page frame 2		<b>2</b>	2	2	<b>1</b>							
	page frame 3			<b>3</b>	3	3							
control states (age per page frame)	page frame 1	0	1	2	0	1							
	page frame 2	>	0	1	2	0							
	page frame 3	>	>	0	1	2							



# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4	4						
	page frame 2		<b>2</b>	2	2	<b>1</b>	1						
	page frame 3			<b>3</b>	3	3	<b>2</b>						
control states (age per page frame)	page frame 1	0	1	2	0	1	2						
	page frame 2	>	0	1	2	0	1						
	page frame 3	>	>	0	1	2	0						

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4	4	<b>5</b>					
	page frame 2		<b>2</b>	2	2	<b>1</b>	1	1					
	page frame 3			<b>3</b>	3	3	<b>2</b>	2					
control states (age per page frame)	page frame 1	0	1	2	0	1	2	0					
	page frame 2	>	0	1	2	0	1	2					
	page frame 3	>	>	0	1	2	0	1					

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4	4	<b>5</b>	5				
	page frame 2		<b>2</b>	2	2	<b>1</b>	1	1	1				
	page frame 3			<b>3</b>	3	3	<b>2</b>	2	2				
control states (age per page frame)	page frame 1	0	1	2	0	1	2	0	1				
	page frame 2	>	0	1	2	0	1	2	3				
	page frame 3	>	>	0	1	2	0	1	2				

# First-In, First-Out (FIFO)

- oldest page is replaced
- necessary states:
  - age (time of swap-in) for each page frame
- sequence of replacement

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4	4	<b>5</b>	5	5	5	5	5
	page frame 2		<b>2</b>	2	2	<b>1</b>	1	1	1	1	<b>3</b>	3	3
	page frame 3			<b>3</b>	3	3	<b>2</b>	2	2	2	2	<b>4</b>	4
control states (age per page frame)	page frame 1	0	1	2	0	1	2	0	1	2	3	4	5
	page frame 2	>	0	1	2	0	1	2	3	4	0	1	2
	page frame 3	>	>	0	1	2	0	1	2	3	4	0	1

- 9 swap-in operations

# First-In, First-Out (FIFO)

- larger main memory with 4 page frames
  - 10 swap-in operations
- FIFO anomaly (Bélády's anomaly, 1969)

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	1	1	1	<b>5</b>	5	5	5	<b>4</b>	4
	page frame 2		<b>2</b>	2	2	2	2	2	<b>1</b>	1	1	1	<b>5</b>
	page frame 3			<b>3</b>	3	3	3	3	3	<b>2</b>	2	2	2
	page frame 4				<b>4</b>	4	4	4	4	4	<b>3</b>	3	3
control states (age per page frame)	page frame 1	0	1	2	3	4	5	0	1	2	3	0	1
	page frame 2	>	0	1	2	3	4	5	0	1	2	3	0
	page frame 3	>	>	0	1	2	3	4	5	0	1	2	3
	page frame 4	>	>	>	0	1	2	3	4	5	0	1	2

# Optimal Replacement Strategy (OPT)

- longest forward distance (LFD, dt. Vorwärtsabstand)
  - time interval until next access to the corresponding page
- strategy OPT (or MIN) is optimal (with fixed number of page frames): minimum number of swap-in/replacements (here: 7)
  - always replace the page with the longest forward distance

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	1	1	1	1	1	1	<b>3</b>	<b>4</b>	4
	page frame 2		<b>2</b>	2	2	2	2	2	2	2	2	2	2
	page frame 3			<b>3</b>	<b>4</b>	4	4	<b>5</b>	5	5	5	5	5
control states (LFD)	page frame 1	4	3	2	1	3	2	1	>	>	>	>	>
	page frame 2	>	4	3	2	1	3	2	1	>	>	>	>
	page frame 3	>	>	7	7	6	5	5	4	3	2	1	>

# Optimal Replacement Strategy (OPT)

- increase the size of the main memory (4 page frames):  
6 swap-in operations
  - no anomaly

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	1	1	1	1	1	1	1	<b>4</b>	4
	page frame 2		<b>2</b>	2	2	2	2	2	2	2	2	2	2
	page frame 3			<b>3</b>	3	3	3	3	3	3	3	3	3
	page frame 4				<b>4</b>	4	4	<b>5</b>	5	5	5	5	5
control states (LFD)	page frame 1	4	3	2	1	3	2	1	>	>	>	>	>
	page frame 2	>	4	3	2	1	3	2	1	>	>	>	>
	page frame 3	>	>	7	6	5	4	3	2	1	>	>	>
	page frame 4	>	>	>	7	6	5	5	4	3	2	1	>

# Optimal Replacement Strategy – Discussion

- implementation of OPT practically impossible
  - reference sequence would have to be known in advance (oracle)
  - OPT is only useful for comparing strategies
- research for strategies that are as close to OPT as possible
  - for example: **Least Recently Used** (LRU)



# Least Recently Used (LRU)

- longest backward distance
  - time interval since the last access to the corresponding page
- LRU strategy (10 swap-in operations)
  - always replace the page that has not been used for the longest time

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4	4	<b>5</b>	5	5	<b>3</b>	3	3
	page frame 2		<b>2</b>	2	2	<b>1</b>	1	1	1	1	1	<b>4</b>	4
	page frame 3			<b>3</b>	3	3	<b>2</b>	2	2	2	2	2	<b>5</b>
control states (oldest access time)	page frame 1	0	1	2	0	1	2	0	1	2	0	1	2
	page frame 2	>	0	1	2	0	1	2	0	1	2	0	1
	page frame 3	>	>	0	1	2	0	1	2	0	1	2	0

# Least Recently Used (LRU)

- increase the size of the main memory (4 page frames)
- 8 swap-in operations

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	1	1	1	1	1	1	1	1	<b>5</b>
	page frame 2		<b>2</b>	2	2	2	2	2	2	2	2	2	2
	page frame 3			<b>3</b>	3	3	3	<b>5</b>	5	5	5	<b>4</b>	4
	page frame 4				<b>4</b>	4	4	4	4	4	<b>3</b>	3	3
control states (oldest access time)	page frame 1	0	1	2	3	0	1	2	0	1	2	3	0
	page frame 2	>	0	1	2	3	0	1	2	0	1	2	3
	page frame 3	>	>	0	1	2	3	0	1	2	3	0	1
	page frame 4	>	>	>	0	1	2	3	4	5	0	1	2

# Least Recently Used (LRU) – Discussion

## ■ no anomaly

- in general, there is a class of algorithms (stack algorithms) for which no anomaly occurs:
  - with stack algorithms, if there are  $k$  page frames at any time, a subset of the pages is swapped-in that would also be swapped-in if there were  $k+1$  page frames at the same time
  - LRU: there are always the last  $k$  used pages swapped-in
  - OPT: the  $k$  already used pages are swapped-in, which will be accessed next

## ■ problem

- implementation of LRU only possible with hardware support
- every memory access must be considered

# Least Recently Used (LRU) – Discussion

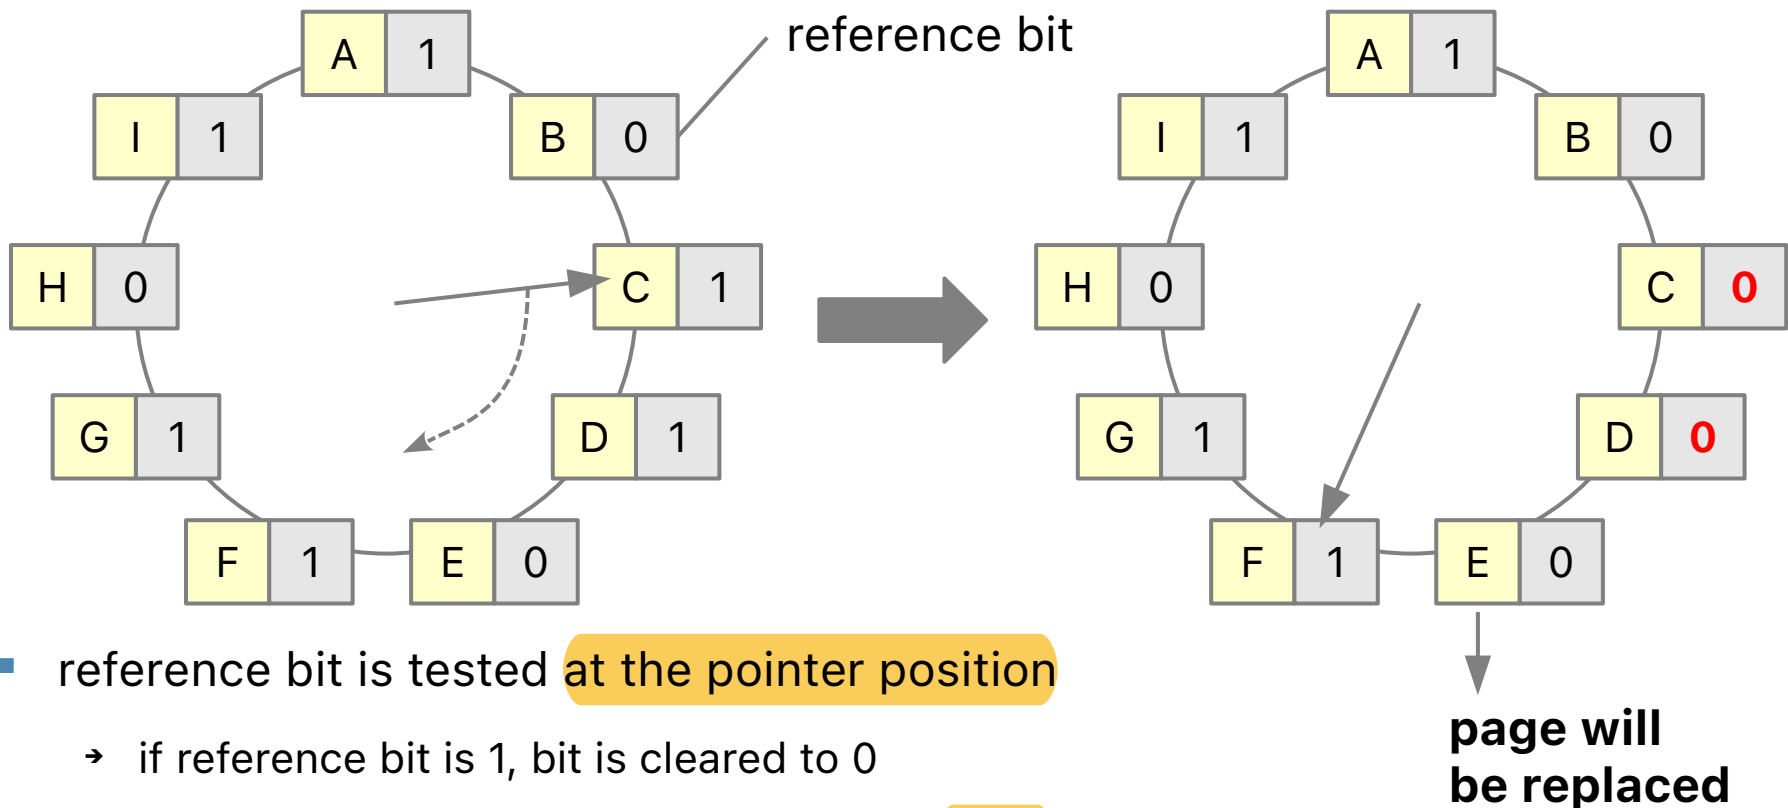
- naive idea: hardware support through counter
  - CPU has a counter that is incremented with each memory access
  - with each memory access: current counter value is written into the respective page descriptor
  - evict the page with the smallest counter reading (→ search!)
- **complex implementation**
  - many additional memory accesses
  - high memory requirements
  - minimum search in page-fault handling

# Second Chance (SC)

- core idea: use of reference bits
  - reference bit in page descriptor is automatically set by hardware when page is accessed
    - easy to implement
    - reduced number of additional memory accesses
    - supported by modern processors and MMUs (e.g., x86: access bit)
- goal: approximation of LRU
  - reference bit is initially set to 1 for a freshly swapped-in page
  - when a victim page is searched, reference bits in the page table are inspected in turn
    - if inspected reference bit is 1, it is set to 0 (second chance)
    - else: if inspected reference bit is 0, the page is chosen as victim and swapped-out

# Second Chance (SC)

- implementation with a circulating pointer (clock)



- reference bit is tested at the pointer position
  - if reference bit is 1, bit is cleared to 0
  - if reference bit is 0, replaceable page is found
  - pointer is incremented; if no victim page is found: repeat
- if all reference bits are set to 1: Second Chance degrades to FIFO

# Second Chance (SC)

- modus operandi with 3 page frames
- 9 swap-in operations

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	<b>4</b>	4	4	<b>5</b>	5	5	5	5	5
	page frame 2		<b>2</b>	2	2	<b>1</b>	1	1	1	1	<b>3</b>	3	3
	page frame 3			<b>3</b>	3	3	<b>2</b>	2	2	2	2	<b>4</b>	4
control states (reference bits)	page frame 1	1	1	<b>1</b>	1	1	<b>1</b>	1	1	1	0	<b>0</b>	<b>1</b>
	page frame 2	<b>0</b>	1	1	<b>0</b>	1	1	<b>0</b>	<b>1</b>	<b>1</b>	1	1	1
	page frame 3	0	<b>0</b>	1	0	<b>0</b>	1	0	0	1	<b>0</b>	1	1
	circulating pointer	<b>2</b>	<b>3</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>1</b>	<b>2</b>	<b>2</b>	<b>2</b>	<b>3</b>	<b>1</b>	<b>1</b>

不用替换时，指针不变，  
但是更新ref bit

# Second Chance (SC)

- increase the size of the main memory (4 page frames)
- 10 swap-in operations

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
main memory	page frame 1	<b>1</b>	1	1	1	1	1	<b>5</b>	5	5	5	<b>4</b>	4
	page frame 2		<b>2</b>	2	2	2	2	2	<b>1</b>	1	1	1	<b>5</b>
	page frame 3			<b>3</b>	3	3	3	3	3	<b>2</b>	2	2	2
	page frame 4				<b>4</b>	4	4	4	4	4	<b>3</b>	3	3
control states (reference bits)	page frame 1	1	1	1	<b>1</b>	<b>1</b>	<b>1</b>	1	1	1	<b>1</b>	1	1
	page frame 2	<b>0</b>	1	1	1	1	1	<b>0</b>	1	1	1	<b>0</b>	1
	page frame 3	0	<b>0</b>	1	1	1	1	0	<b>0</b>	1	1	0	<b>0</b>
	page frame 4	0	0	<b>0</b>	1	1	1	0	0	<b>0</b>	1	0	0
	circulating pointer	<b>2</b>	<b>3</b>	<b>4</b>	<b>1</b>	<b>1</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>	<b>1</b>	<b>2</b>	<b>3</b>

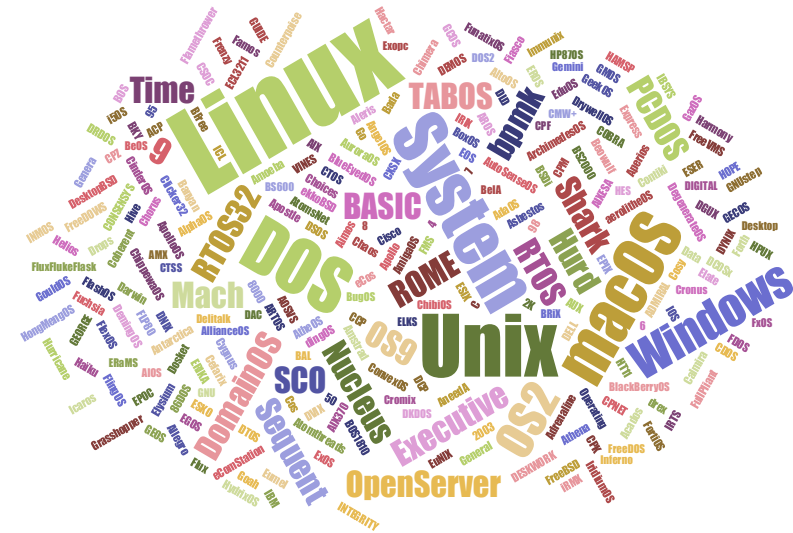


# Second Chance (SC)

- second chance may lead to FIFO anomaly
  - if all reference bits are set, the decision is made according to FIFO
- however, normally Second Chance comes close to LRU
- enhanced Second Chance variant
  - modification bit can be considered additionally (dirty bit)
  - four classes (reference bit, modification bit):
    - (0,0): neither recently used nor modified
    - (0,1): not recently used but modified
    - (1,0): recently used and unmodified
    - (1,1): recently used and modified
  - search for the lowest class

# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ Demand Paging
- ▶ Page Replacement Algorithms
  - ▶ First In, First Out (FIFO)
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ Summary and Outlook



# Page Frame Allocation

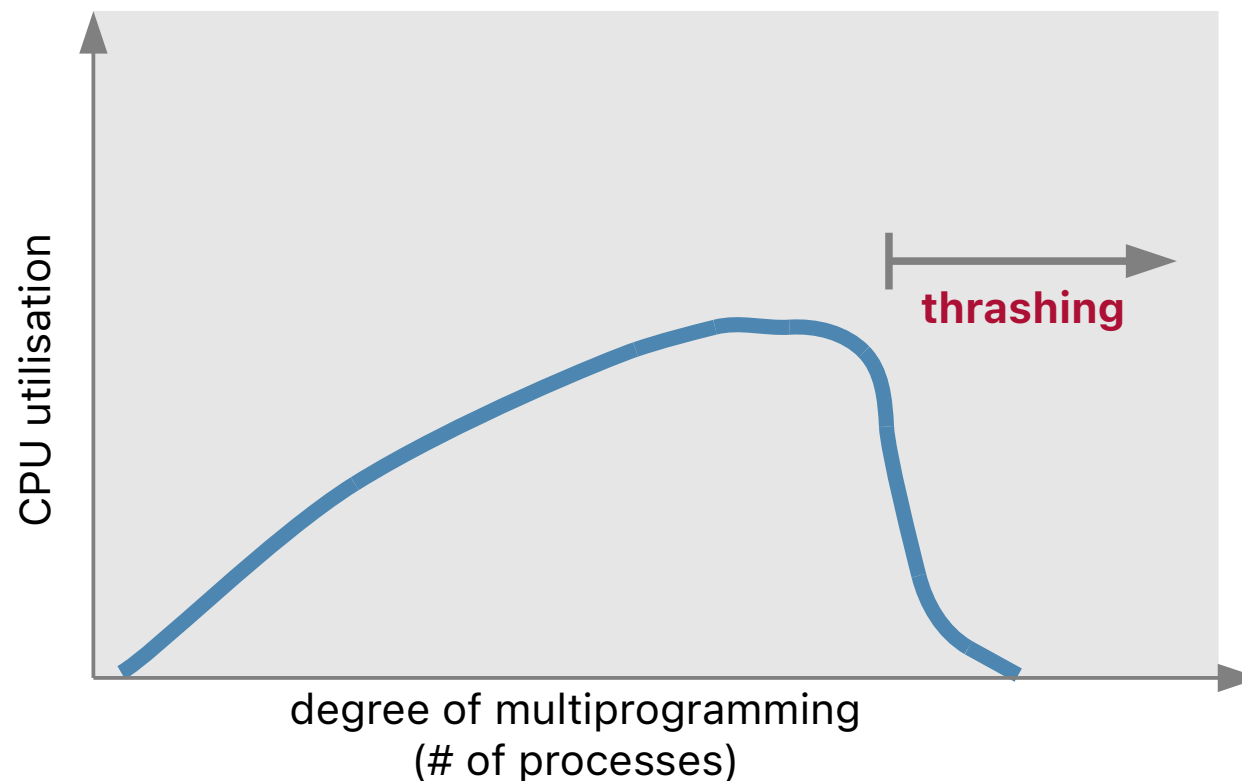
- problem: distribution of available page frames to processes
  - how many swapped-in pages should a process possess?
    - ➔ maximum: limited by size of available memory (number of page frames)
    - ➔ minimum: depends on the processor architecture
      - at least the number of pages necessary which in theory are required for executing a single instruction
      - for example: 2 pages for instructions, 4 pages for the data
- homogenous allocation
  - number of processes determines how many page frames each process can possess
- size-based allocation
  - process size influences page frame allocation

# Page Frame Allocation

- global and local page frame allocation
  - **local**: processes always replace their own pages
    - page fault handling: responsibility of the individual process, only
  - **global**: processes can also displace pages of other processes
    - better efficiency, as unused pages can be used by other processes

# Thrashing

- thrashing (dt. **Seitenflattern**) happens when swapped-out pages are immediately addressed again
  - process spends more time waiting for page faults handled than actually executing its own code



# Thrashing

- causes

- process operates with a number of page frames close to the necessary minimum
- too many processes in the system at the same time
- poor replacement strategy

➔ local page allocation resolves thrashing between processes

➔ allocation of a sufficiently large number of page frames resolves thrashing within the pages of a process

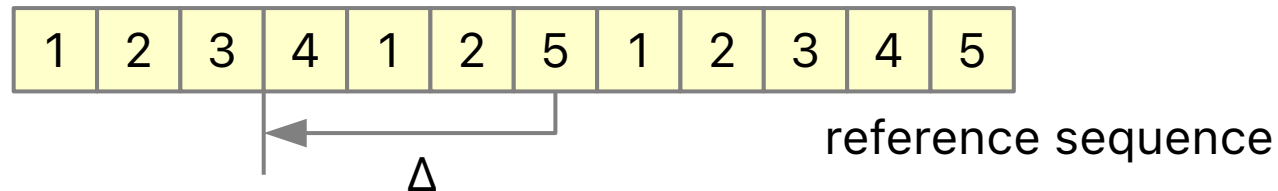
- limitation of the number of processes

# Thrashing – Solution 1

- solution 1: swap-out processes
- inactive processes do not need page frames
  - page frames are distributed to a smaller number of processes
  - needs to be interlinked with scheduling
    - prevent starvation of processes
    - achieve good response times (low latency)

# Thrashing – Solution 2

- solution 2: working set model
- working set: amount of pages that a process really needs
  - can only be approximated, as usually not predictable
- approach by looking at the last  $\Delta$  pages that were referenced
  - appropriate choice of  $\Delta$ 
    - too large: overlapping of local access patterns
    - too small: work set does not contain all necessary pages





# Working Set Model

- example: working set model for different  $\Delta$

reference sequence		1	2	3	4	1	2	5	1	2	3	4	5
$\Delta = 3$	page 1	X	X	X		X	X	X	X	X	X		
	page 2		X	X	X		X	X	X	X	X	X	
	page 3			X	X	X					X	X	X
	page 4				X	X	X					X	X
	page 5							X	X	X			X
$\Delta = 4$	page 1	X	X	X	X	X	X	X	X	X	X	X	
	page 2		X	X	X	X	X	X	X	X	X	X	X
	page 3			X	X	X	X				X	X	X
	page 4				X	X	X	X				X	X
	page 5							X	X	X	X		X

# Working Set Model with Timer

- implementation challenges similar to LRU: track how long pages have not been referenced
- algorithm with timer interrupt
  - regular interrupts update the page age information by means of the reference bit:
    - if reference bit is set (page was used), age is set to zero;
    - otherwise: age is increased
    - only the pages of the currently running process are "aged"
  - pages with age  $> \Delta$  are no longer part of the working set of the respective process
- impractical as to run-time overhead

# Thrashing – Solution 3

- solution 3: react on rate of page faults
- thrashing can be more easily avoided by directly controlling the page fault rate
  - page-fault logging per process
  - page-fault rate  $<$  threshold: decrease number of page frames
  - page-fault rate  $>$  threshold: increase number of page frames

# Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Motivation
  - ▶ Locality of Reference
  - ▶ Virtual Memory
- ▶ Demand Paging
- ▶ Page Replacement Algorithms
  - ▶ First In, First Out (FIFO)
  - ▶ Least Recently Used (LRU)
  - ▶ Second Chance (SC)
- ▶ Page Frame Allocation
  - ▶ Thrashing
  - ▶ Working Set Model
- ▶ **Summary and Outlook**



# ▶▶ Summary and Outlook

## ■ summary

- virtual memory enables the use of large logical address spaces despite limitations of the main memory
- but: comfort comes at a price
  - increased complexity and effort at the hardware level
  - complex algorithms in the operating system → runtime overhead
  - strategy selection depends on the usage pattern
- more simple (special-purpose) operating systems that do not necessarily need this "luxury" are better off without it

## ■ outlook: filesystems

- abstraction methods to represent data blocks as files and directories
- data management and storage, error handling and recovery

# References and Acknowledgments

## Lecture

- ▶ Systemnahe Programmierung in C (SPiC), Betriebssysteme (Jürgen Kleinöder, Wolfgang Schröder-Preikschat)
- ▶ Betriebssysteme und Rechnernetze (Olaf Spinczyk, Embedded Software Systems Group, Universität Osnabrück)

## Teaching Books and Reference Book

- [1] Avi Silberschatz, Peter Baer Galvin, Greg Gagne: *Operating System Concepts*, John Wiley & Sons, 2018.
- [2] Andrew Tanenbaum, Herbert Bos: *Modern Operating Systems*, Pearson, 2015.
- [3] Wolfgang Schröder-Preikschat: *Grundlage von Betriebssystemen – Sachwortverzeichnis*, 2023.  
<https://www4.cs.fau.de/~wosch/glossar.pdf>