

Operating Systems

Timo Hönig

Bochum Operating Systems and System Software (BOSS)
Ruhr University Bochum (RUB)

XI. Inter-Process Communication

June 28, 2023 (Summer Term 2023)



RUHR
UNIVERSITÄT
BOCHUM

RUB

www.informatik.rub.de
Chair of Operating Systems and System Software

Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Principles of Inter-Process Communication
- ▶ Local Inter-Process Communication with UNIX-style Operating Systems
 - ▶ Signals
 - ▶ Pipes
 - ▶ Message Queues
- ▶ Inter-Process Communication Across Systems
 - ▶ Sockets
 - ▶ Remote Procedure Calls
- Summary and Outlook




Organizational Matters

■ lecture

- Wednesday, 10:15 – 11:45
- format: synchronous, **hybrid**
 - in presence (Room HID, Building ID)
 - online lecture (Zoom)
- **exam:** August 7, 2023 (first appointment)
September 25, 2023 (retest appointment)
- **evaluation is live!**

<https://tinyurl.com/25twb8qr>

- manage course material, asynchronous communication: Moodle
- <https://moodle.ruhr-uni-bochum.de/course/view.php?id=50698>



Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Principles of Inter-Process Communication
- ▶ Local Inter-Process Communication with UNIX-style Operating Systems
 - ▶ Signals
 - ▶ Pipes
 - ▶ Message Queues
- ▶ Inter-Process Communication Across Systems
 - ▶ Sockets
 - ▶ Remote Procedure Calls
- Summary and Outlook



Principles of Inter-Process Communication

- processes interact
 - wait for each other: **synchronisation**
 - share and exchange data: **communication**
- waiting mechanisms
 - are necessary for controlled communication
 - potentially lead to **deadlock** situations
- up to now: data **exchange** has only had minor consideration
 - lightweight and featherweight processes
 - in the same address space
 - using different address spaces
 - transfer of arbitrary types and amounts of data

Inter-Process Communication (IPC)

- multiple processes work jointly on **one** task
 - **simultaneous use** of available information **by several processes**
 - increase performance/reduce processing time by parallelisation
 - hiding of processing times by execution "in the background"
- communication through **shared memory**
 - data exchange: concurrent writing to (or reading from) a shared memory
 - attention must be paid to synchronisation
- **focus:** communication through **messages**
 - messages are exchanged between processes
 - shared memory is **not** mandatory

Message-based Communication

- communication primitives (typical signature):

```
send (destination, message)
```

```
receive (source, message)
```

- there are differences in:
 - synchronisation method
 - addressing mode
 - additionally: several properties

Synchronisation for Message-based IPC

- **synchronisation on send / receive**
 - **synchronous message exchange** (also „rendezvous“)
 - recipient blocks until the message is received
 - sender blocks until the arrival of the message is confirmed
 - **asynchronous message exchange**
 - sender passes the message to the operating system and then resumes operation
 - blocking on both sides optional
 - buffering always required
- frequently found:
 - asynchronous message exchange with potentially blocking sending and receiving

Addressing Modes for Message-based IPC

- **direct addressing**

- process ID (e.g., signals)
- communication endpoint of a process (e.g., port, socket)

- **indirect addressing**

- channels (e.g., pipes)
- mail boxes, message buffer (e.g., message queue)

- additional dimension: **addressing of groups**

- **unicast** – send message to exactly one process
- **multicast** – send message to a selection of processes
- **broadcast** – send message to all processes

Additionally Properties of Message-based IPC

- **message format**

- stream oriented vs. message oriented
- fixed length vs. variable length
- typed vs. untyped

- **transmission**

- unidirectional vs. bidirectional (i.e., half-duplex, full-duplex)
- reliable vs. unreliable
- order of messages is kept vs. is not kept

Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Principles of Inter-Process Communication
- ▶ **Local Inter-Process Communication with UNIX-style Operating Systems**
 - ▶ Signals
 - ▶ Pipes
 - ▶ Message Queues
- ▶ Inter-Process Communication Across Systems
 - ▶ Sockets
 - ▶ Remote Procedure Calls
- Summary and Outlook



Signals


- signals are interrupts recreated in software
 - similar to interrupts of a processor through I/O devices
 - **minimal form** of inter-process communication (transmission of the signal number, **no payload**)
- sender:
 - processes - with the help of the system call **kill(2)**
 - operating system - when certain events occur
- receiver process performs signal handling:
 - ignore
 - terminate process
 - call a signal handler function
 - ➔ after handling the signal, the process continues at the interrupted position

Signals

- with the help of signals, processes can be informed about exceptional situations (c.f., hardware interrupts)
- examples:
 - **SIGINT** abort process (e.g., Ctrl-C)
 - **SIGSTOP** stop process (e.g., Ctrl-Z)
 - **SIGCHLD** child process terminated
 - **SIGSEGV** memory protection violation of the process
 - **SIGKILL** process is killed
- the default signal handling (terminate, stop, ...) can be redefined for most signals
 - see **signal(7)**

Signals – Logical View

Hollywood Principle: „*Don't call us, we'll call you.*“




Signals – Technical View

- signal handling always takes place at the transition from kernel to user mode
- What happens when the receiving process...
 1. runs in state **RUNNING** (e.g., segmentation error, bus error)?
 - immediate start of the signal handling routine
 2. does not run but is in state **READY** (e.g., `kill` system call)?
 - the signal is registered in the process control block (PCB)
 - as soon as the process runs again, the signal handling takes place
 3. waits on an I/O event, in state **BLOCKED**?
 - the I/O system call (e.g., `read`) is aborted with `EINTR`
 - the process state is set to **READY**
 - after that: as for 2.
 - if applicable, the interrupted system call is executed again (`SA_RESTART`)

IPC with Pipes

- channel between two communication partners
 - unidirectional
 - buffered (fixed buffer size)
 - reliable
 - stream oriented
- pipe operations: read and write
 - order of the bytes is preserved (byte stream)
 - processes block when the pipe is full (write) or empty (read)



Pipes – Programming

■ unnamed pipes

- create a pipe:

```
int pipe (int pipefd[2])
```

- After a successful call (return value is zero):

- **read from** the pipe via `pipefd[0]` (system call `read`)
- **write to** the pipe via `pipefd[1]` (system call `write`)

- now one just has to pass one end of the pipe to another process (next slide)

■ named pipes

- pipes can also be placed in the file system as special files:

```
int mkfifo (const char *filename, mode_t mode)
```

- standard functions (i.e., `open`, `read`, `write` and `close`) can be used
 - file access rights control who can use the pipe

Pipes – Programming

```
enum { READ=0, WRITE=1 };

int main (int argc, char *argv[]) {
    int res, fd[2];
    if (pipe (fd) == 0) {                                /* create pipe */
        res = fork ();
        if (res > 0) {                                    /* parent process */
            close (fd[READ]);                             /* close reading side */
            dup2 (fd[WRITE], 1);                          /* write stdout to pipe */
            close (fd[WRITE]);                             /* free descriptor */
            execlp (argv[1], argv[1], NULL);               /* execute writer */
        }
        else if (res == 0) {                              /* child process */
            close (fd[WRITE]);                             /* close writing side */
            dup2 (fd[READ], 0);                            /* read from pipe as stdin */
            close (fd[READ]);                              /* free descriptor */
            execlp (argv[2], argv[2], NULL);               /* execute reader */
        }
    }
    /* error handling */
}
```

Pipes – Programming

```
enum { READ=0, WRITE=1 };
```

```
int main (int argc, char *argv[]) {
    int res, fd[2];
    if (pipe (fd) == 0) {
        res = fork ();
        if (res > 0) {
```

**"./connect ls wc" is
equivalent to the shell
command "ls | wc"**

```
/* create pipe */
```

```
/* parent process */
```

```
user@host$ ls
connect  connect.c  execl.c  fork.c  orphan.c  wait.c
user@host$ ./connect ls wc
        6        6       49
```

```
close (fd[WRITE]); /* close writing side */
dup2 (fd[READ], 0); /* read from pipe as stdin */
close (fd[READ]); /* free descriptor */
execlp (argv[2], argv[2], NULL); /* execute reader */
```

```
    }
}
```

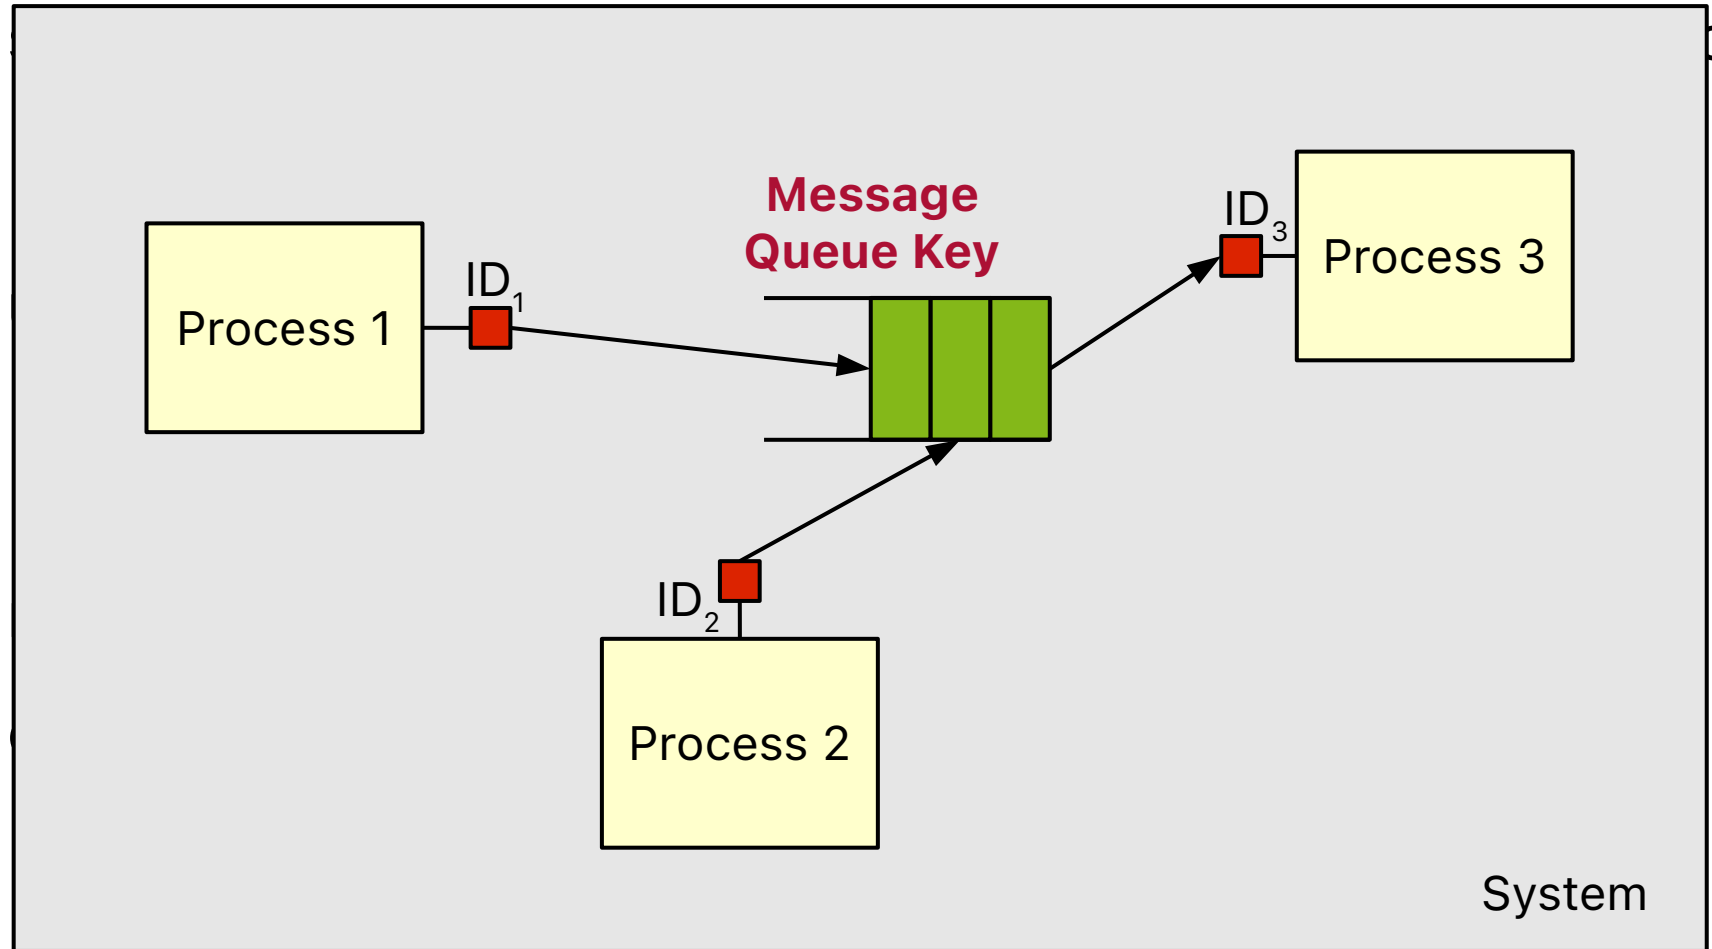
```
/* error handling */
```

```
}
```

IPC with Message Queues

- system-wide unique address (key) serves for identification
 - access rights as to files
 - process local number (`msqid`) is required for **all** operations
- undirected M:N communication
- buffered operation
 - adjustable size per queue
- messages are typed (long)
- operations to send and receive a message
 - blocking vs. non-blocking
 - receive all messages vs. receive a specific message type only

IPC with Message Queues



Message Queues – Programming

- create a message queue and get message queue id:

```
int msgget (key_t key, int msgflg) ;
```

- all communicating processes must know the key
- keys are unique within an operating system instance
- if a key is already assigned, **no** message queue with the same key can be created

- message queues can be created **without** a key (private queues, **key=IPC_PRIVATE**)

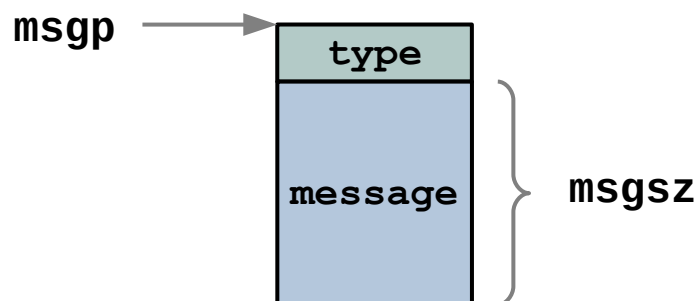
- non-private message queues are persistent
- they must be explicitly deleted (**cmd=IPC_RMID**):

```
int msgctl (int msqid, int cmd  
            struct msqid_ds *buf) ;
```


Message Queues – Programming

- send a message

```
int msgsnd (int msqid, const void *msgp,  
            size_t msgsz, int msgflg);
```



- receive a message

```
int msgrcv (int msqid, void *msgp, size_t msgsz,  
            long msgtype, int msgflg);
```

- `msgtype = 0`: first message
- `msgtype > 0`: first message with given type
- `msgtype < 0`: message with smallest type $\leq |\text{msgtype}|$

Message Queues – Commands

- show all message queues

```
ipcs -q
```

- remove message queue

```
ipcrm -Q <key>
```

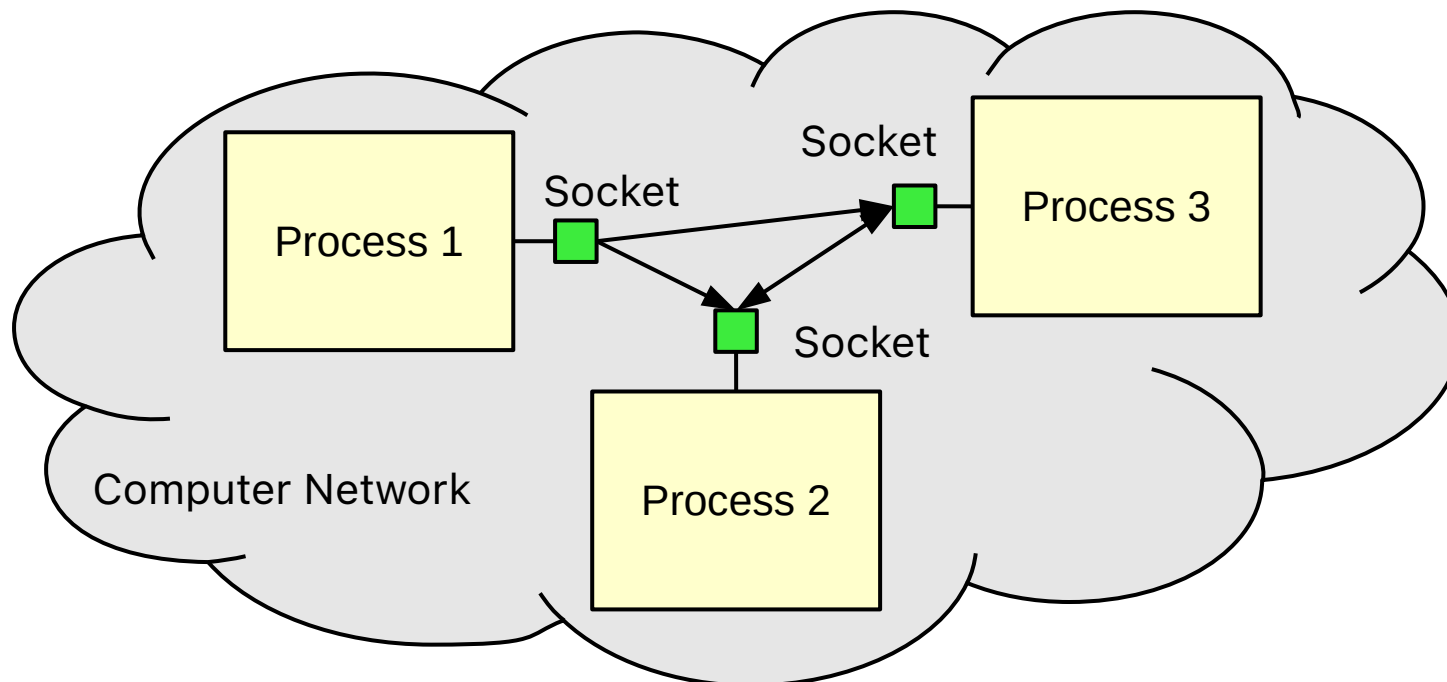
Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Principles of Inter-Process Communication
- ▶ Local Inter-Process Communication with UNIX-style Operating Systems
 - ▶ Signals
 - ▶ Pipes
 - ▶ Message Queues
- ▶ Inter-Process Communication Across Systems
 - ▶ Sockets
 - ▶ Remote Procedure Calls
- Summary and Outlook



Sockets

- sockets are **general communication endpoints** in a computer network
 - bidirectional
 - buffered
- abstracts from details of the communication system
 - characterised by **domain** (protocol family), **type** and **protocol**



Socket Domains

- UNIX domain socket
 - UNIX **domain sockets** behave like **bidirectional pipes**
 - creation as a special file in the file system is possible
- Internet domain socket
 - designed for inter-computer communication using Internet protocols
 - Appletalk Domain, DECnet Domain, ...
- domains determine possible protocols
 - for example, Internet domain: TCP/IP or UDP/IP
- domains determine the address family
 - example: Internet domain - IP address and port number

Socket Types and Protocols

- the major types of sockets:
 - **SOCK_STREAM**: stream-oriented, connection-oriented, reliable (TCP)
 - **SOCK_DGRAM**: message-oriented, unreliable (UDP)
 - Internet domain protocols:
 - TCP/IP protocol
 - stream- and connection-oriented, reliable
 - UDP/IP protocol
 - message-oriented, connectionless, unreliable
 - messages can be lost or duplicated
 - sequence can get out of order
 - packet boundaries are preserved (datagram protocol)
- protocol specification is often redundant

Socket Programming

- creating sockets
 - generate a socket with system call **socket** (return value is a file descriptor):

```
int socket (int domain, int type, int protocol);
```

- address binding
 - sockets are generated without address
 - binding to an address is done by:

```
int bind (int sockfd,  
          const struct sockaddr_in *addr,  
          socklen_t addrlen);
```

- **struct sockaddr_in** (for Internet address family) contains:

sin_family: AF_INET

sin_port: 16 bit port number

sin_addr: struct with IP address (e.g., 192.168.2.1)

note: for IPv6 there are
sockaddr_in6 and AF_INET6

Socket Programming

■ datagram sockets

- packet-oriented
- no connection setup necessary
- **sending** a datagram:

```
ssize_t sendto (int sockfd, const void *buf,  
               size_t len, int flags,  
               const struct sockaddr *dest_addr,  
               socklen_t addrlen);
```

- **receiving** a datagram:

```
ssize_t recvfrom (int sockfd, void *buf,  
                 size_t len, int flags,  
                 struct sockaddr *src_addr,  
                 socklen_t *addrlen);
```


Socket Programming

- **stream sockets**

- stream-oriented
- connection setup necessary
- client (user, user program) wants to establish a **communication connection** to a server (service provider)

- client: connection setup for stream-oriented sockets

- connecting the socket with:

```
int connect (int sockfd,  
             const struct sockaddr *addr,  
             socklen_t addrlen);
```

- send and receive with **write** and **read** (or **send** and **recv**)
- terminate the connection with **close** (closes the socket)

Socket Programming

- server: accepts requests/orders
 - binds socket to an address (otherwise not reachable)
 - prepares socket for connection requests with:

```
int listen (int sockfd, int backlog);
```

- accepts individual connection requests with:

```
int accept (int sockfd, struct sockaddr *addr,  
            socklen_t *addrlen);
```

- returns a **new socket** connected to the client
 - blocks if there is **no** connection request
- reads data with **read** and executes the implemented service
- sends the result with **write** back to the sender
- closes the new socket with **close**

Socket Programming

```
#define PORT 4711
#define MAXREQ (4096*1024)

char buffer[MAXREQ], body[MAXREQ], msg[MAXREQ];

void error(const char *msg) { perror(msg); exit(1); }

int main() {
    int sockfd, newsockfd;
    socklen_t clilen;
    struct sockaddr_in serv_addr, cli_addr;
    int n;
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
    if (sockfd < 0) error("ERROR opening socket");
    bzero((char *) &serv_addr, sizeof(serv_addr));
    serv_addr.sin_family = AF_INET;
    serv_addr.sin_addr.s_addr = INADDR_ANY;
    serv_addr.sin_port = htons(PORT);
    if (bind(sockfd, (struct sockaddr *) &serv_addr, sizeof(serv_addr)) < 0)
        error("ERROR on binding");
    listen(sockfd, 5);
    ...
}
```

socket is created
and bound to an
address

create the new socket with socket

Socket Programming

```
...
while (1) {
    cliilen = sizeof(cli_addr);
    newsockfd = accept (sockfd, (struct sockaddr *) &cli_addr, &cliilen);
    if (newsockfd < 0) error("ERROR on accept");
    bzero(buffer, sizeof(buffer));
    n = read (newsockfd, buffer, sizeof(buffer)-1);
    if (n < 0) error("ERROR reading from socket");
    snprintf (body, sizeof (body),
        "<html>\n<body>\n"
        "<h1>Hi Browser</h1>\nYour request was: ...\n"
        "<pre>%s</pre>\n"
        "</body>\n</html>\n", buffer);
    snprintf (msg, sizeof (msg),
        "HTTP/1.0 200 OK\n"
        "Content-Type: text/html\n"
        "Content-Length: %d\n\n%s", strlen (body), body);
    n = write (newsockfd, msg, strlen(msg));
    if (n < 0) error("ERROR writing to socket");
    close (newsockfd);
}
}
```

accepting a new connection


```
read HTTP request
```

```
generate reply and
send back to client
```

close connection

Remote Procedure Call (RPC)

- function call across process boundaries
(dt. Fernaufruf)
 - high level of abstraction
 - remote procedure calls are rarely offered directly by the system;
needs mapping to other forms of communication (e.g., to messages)
 - mapping to multiple messages
 - ➔ job message transports call intention and parameters
 - ➔ result message transports results of the call



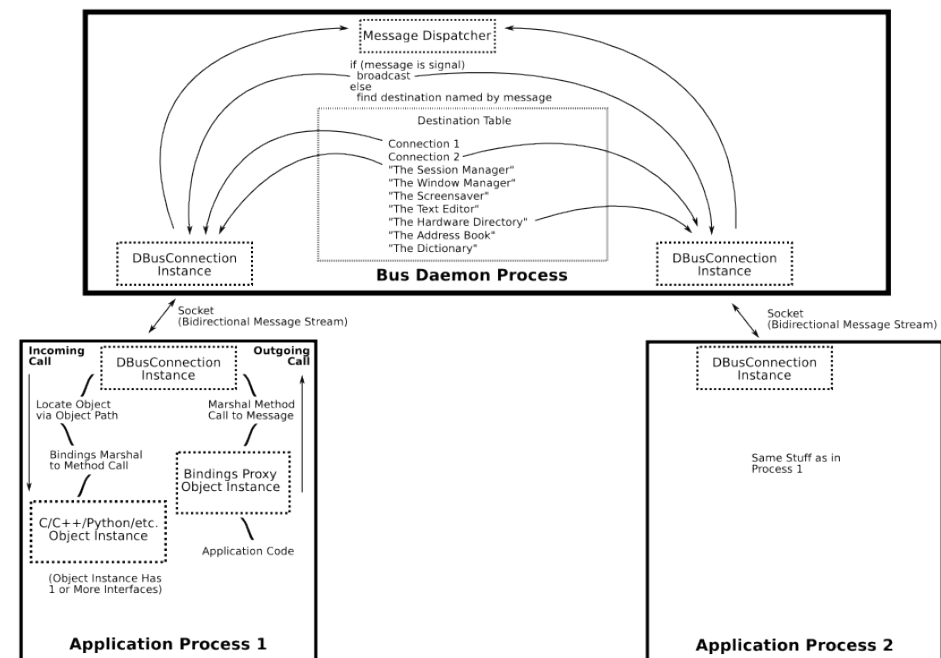
- Examples: NFS, D-Bus

Remote Procedure Call (RPC) – Discussion

- flexible means to dynamically provide functionalities of one process to others
- RPCs latencies
 - using RPCs can lead to high overheads
 - RPC overheads apply to both, on-system IPC and across-system IPC

■ Example: D-Bus

- high number of context switches
- comfort comes at the expense of performance



Agenda

- ▶ Recap
- ▶ Organizational Matters
- ▶ Principles of Inter-Process Communication
- ▶ Local Inter-Process Communication with UNIX-style Operating Systems
 - ▶ Signals
 - ▶ Pipes
 - ▶ Message Queues
- ▶ Inter-Process Communication Across Systems
 - ▶ Sockets
 - ▶ Remote Procedure Calls
- Summary and Outlook



▶▶ Summary and Outlook

■ summary

- types of interprocess communication: **message-based** and **shared memory**
- message-based IPC
 - data is copied
 - works across computer boundaries
- UNIX systems offer various abstractions
 - signals, pipes, sockets, message queues
 - sockets: standardised interface,
implemented by all general-purpose operating systems

■ outlook: multi and many-core systems, systems research

- challenges in operating system design for large number of processing cores
- advanced topics in operating system design, research topics at the chair

References and Acknowledgments

Lecture

- ▶ Systemnahe Programmierung in C (SPiC), Betriebssysteme (Jürgen Kleinöder, Wolfgang Schröder-Preikschat)
- ▶ Betriebssysteme und Rechnernetze (Olaf Spinczyk, Embedded Software Systems Group, Universität Osnabrück)

Teaching Books and Reference Book

- [1] Avi Silberschatz, Peter Baer Galvin, Greg Gagne: *Operating System Concepts*, John Wiley & Sons, 2018.
- [2] Andrew Tanenbaum, Herbert Bos: *Modern Operating Systems*, Pearson, 2015.
- [3] Wolfgang Schröder-Preikschat: *Grundlage von Betriebssystemen – Sachwortverzeichnis*, 2023.
<https://www4.cs.fau.de/~wosch/glossar.pdf>